



Machine Learning Foundations - Part 3

- by *Nagaraju Budigam*

Learning Goals

❑ Regression Metrics

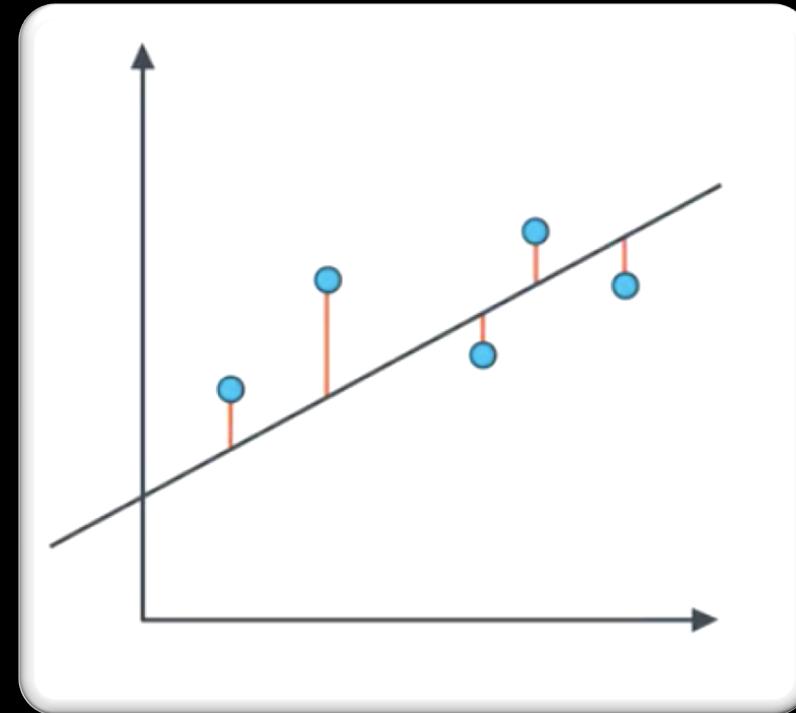
- ✓ Mean Absolute Error
- ✓ Mean Squared Error

❑ Error Detection

- ✓ Underfitting
- ✓ Overfitting
- ✓ Model Complexity Graph
- ✓ Learning Curves
- ✓ K Fold Cross Validation

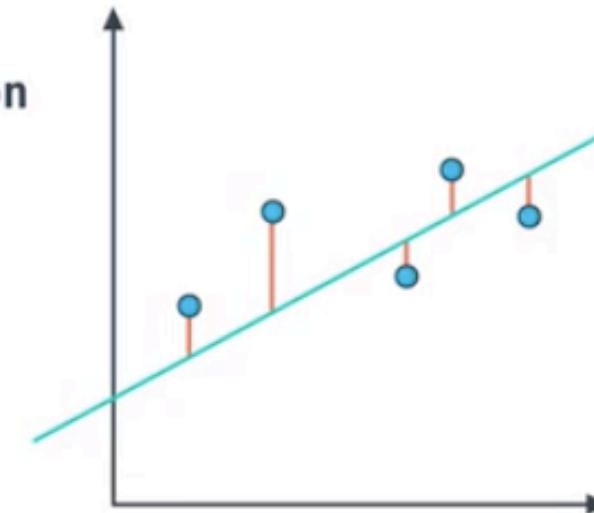
Regression Metrics – Mean Absolute Error

$$\text{MAE} = \frac{1}{n} \sum_{j=1}^n |y_j - \hat{y}_j|$$



Regression Metrics – Mean Absolute Error

```
from sklearn.metrics import mean_absolute_error  
from sklearn.linear_model import LinearRegression  
  
classifier = LinearRegression()  
classifier.fit(X,y)  
  
guesses = classifier.predict(X)  
  
error = mean_absolute_error(y, guesses)
```



Regression Metrics – Mean Absolute Error

❑ Pros

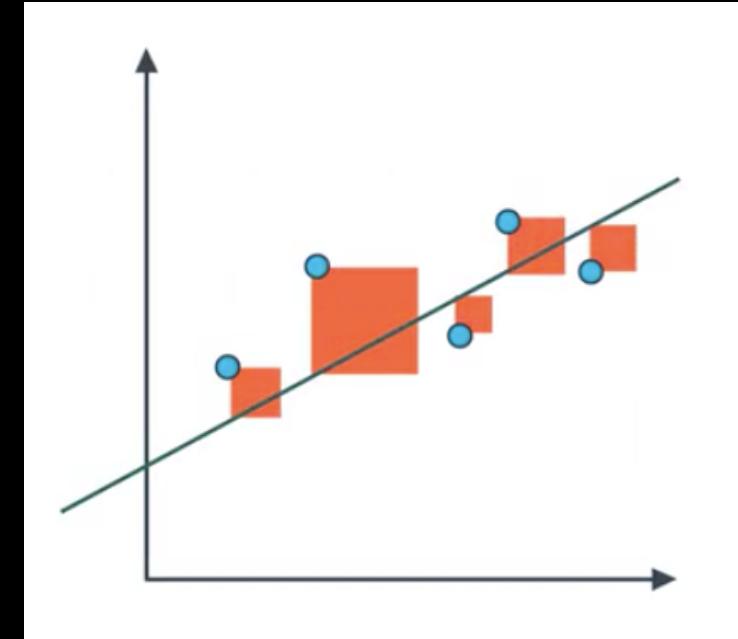
- ✓ Less sensitive to outliers
- ✓ Many Small errors is same as a large error.

❑ Cons

- Absolute Function is not differentiable.
- Not a good option you want to use Gradient Descent as a error metric.

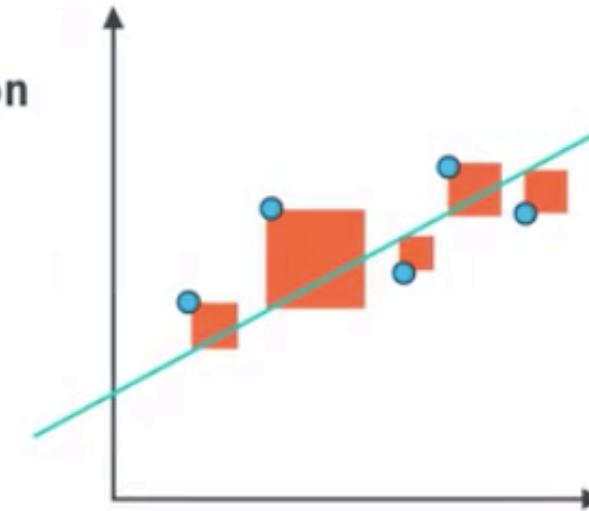
Regression Metrics – Mean Squared Error

$$\text{MSE} = \sqrt{\frac{\sum_{i=1}^n (P_i - O_i)^2}{n}}$$



Regression Metrics – Mean Squared Error

```
from sklearn.metrics import mean_squared_error  
from sklearn.linear_model import LinearRegression  
  
classifier = LinearRegression()  
classifier.fit(X,y)  
  
guesses = classifier.predict(X)  
  
error = mean_squared_error(y, guesses)
```



Regression Metrics – Mean Squared Error

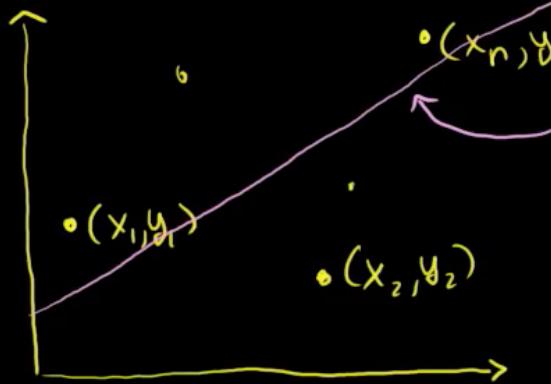
□ Pros

- ✓ Emphasizes extremes than small errors
- ✓ Square Function is differentiable.
- ✓ Good option you want to use Gradient Descent as a error metric.

□ Cons

- Very sensitive to outliers.

Regression Metrics – R2 Score (Coefficient of Determination)



total variation in y :

$$SE_{\bar{y}} = \sqrt{(y_1 - \bar{y})^2 + (y_2 - \bar{y})^2 + \dots + (y_n - \bar{y})^2}$$

$$\frac{SE_{\text{line}}}{SE_{\bar{y}}}$$

How much of total variation
is not described by the regression
line?

$$1 - \frac{SE_{\text{line}}}{SE_{\bar{y}}} = \text{what \% of total variation
is desc by the variation in}$$

Regression Metrics – R2 Score (Coefficient of Determination)



BAD MODEL

The errors should be similar.
R2 score should be close to 0.

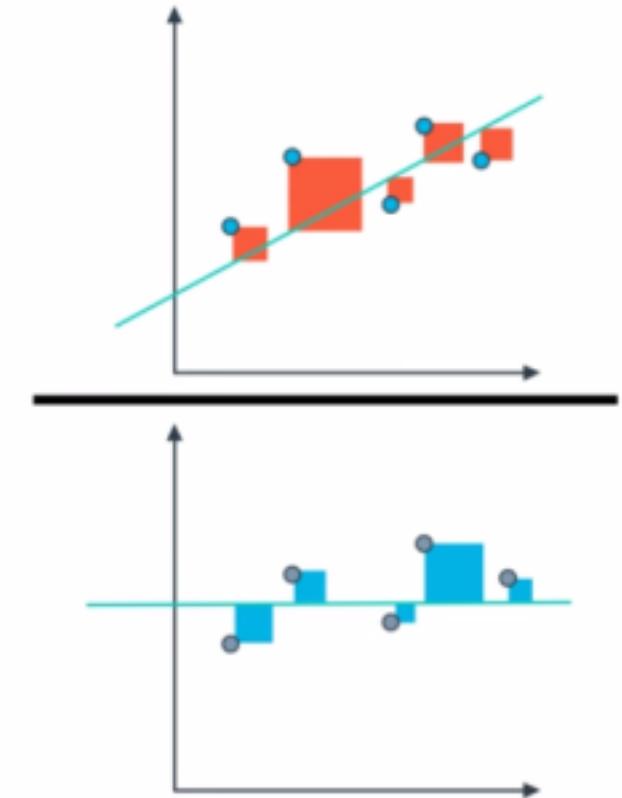


GOOD MODEL

The mean squared error for the linear regression model should be a lot smaller than the mean squared error for the simple model.

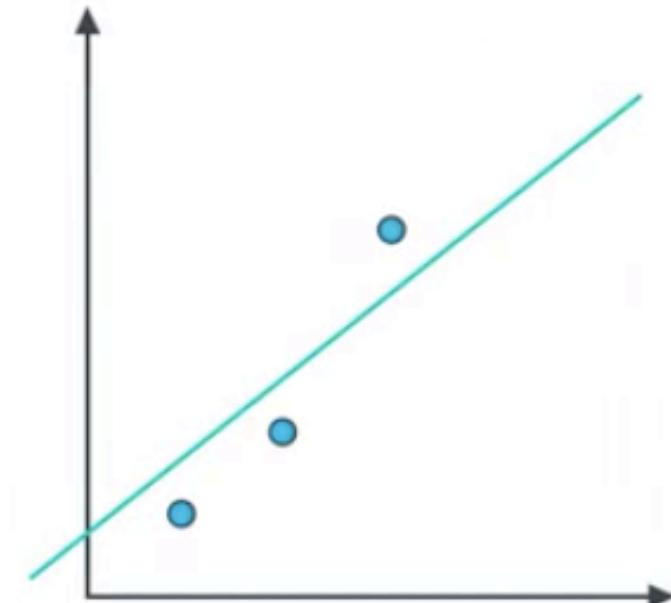
R2 score should be close to 1.

$$R^2 = 1 -$$

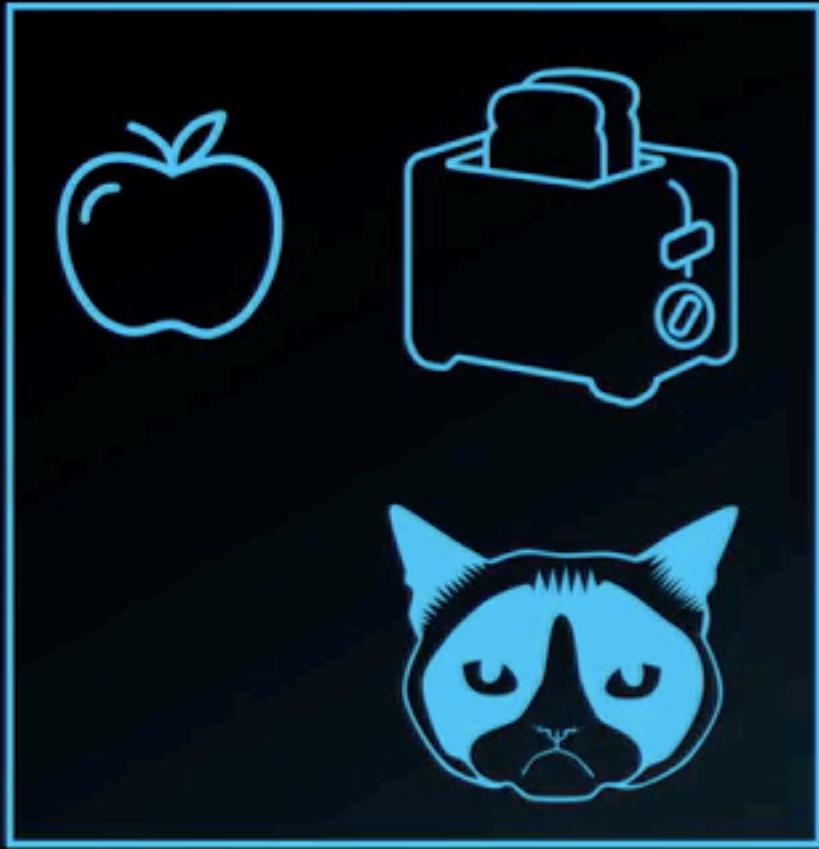


Regression Metrics – R2 Score (Coefficient of Determination)

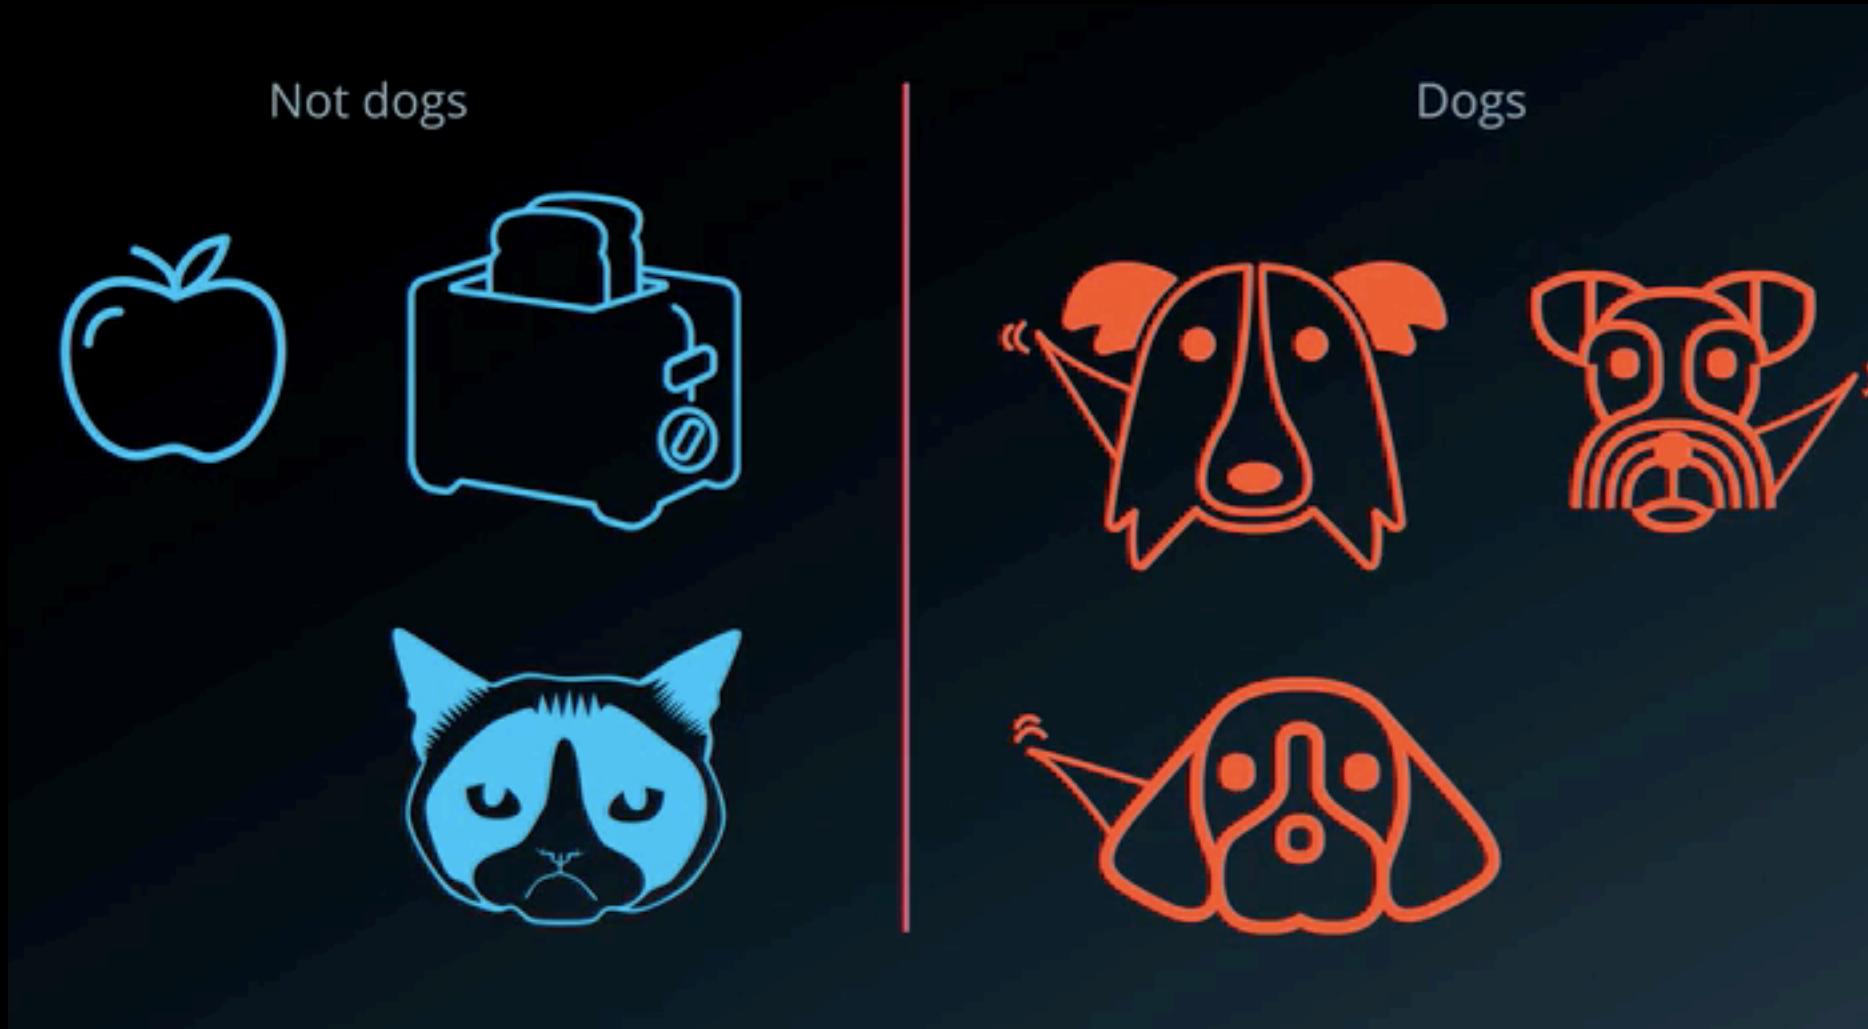
```
from sklearn.metrics import r2_score  
  
y_true = [1, 2, 4]  
y_pred = [1.3, 2.5, 3.7]  
  
r2_score(y_true, y_pred)
```



Error Detection – Types of Errors

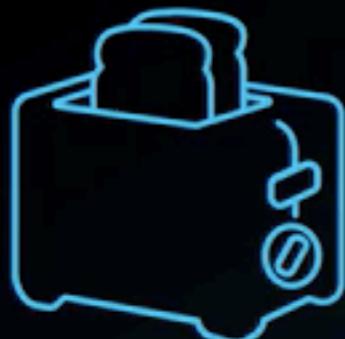


Error Detection – Types of Errors



Error Detection – Types of Errors

Not animals



Animals



Error Detection – Types of Errors

- OVERSIMPLIFIED

Not animals



Animals



Error Detection – Types of Errors - Underfitting

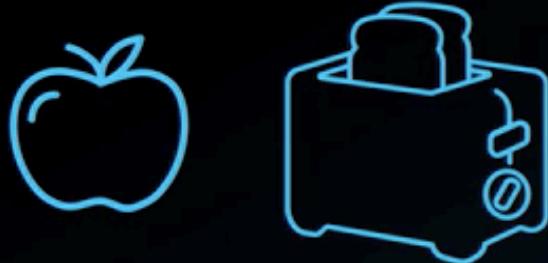
- Does not do well in the training set.
- Error due to bias.



Error Detection – Types of Errors

- OVERCOMPLICATED

No dogs who wag their tail



Dogs that are wagging their tail

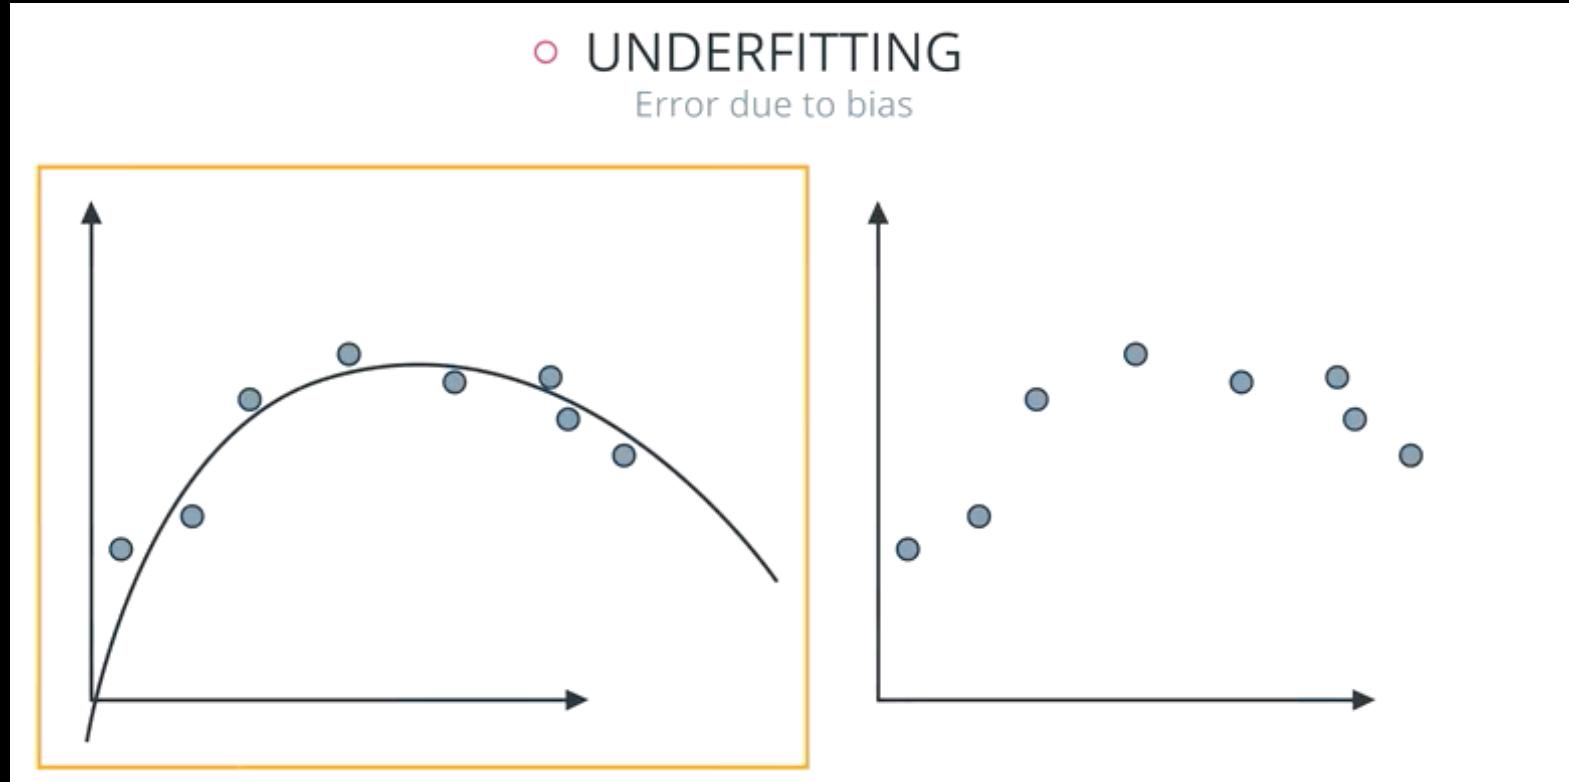


Error Detection – Types of Errors - Overfitting

- Does well in the training set, but it tends to memorize it instead of learning the characteristics of it.
- Error due to variance.



Error Detection – Underfitting

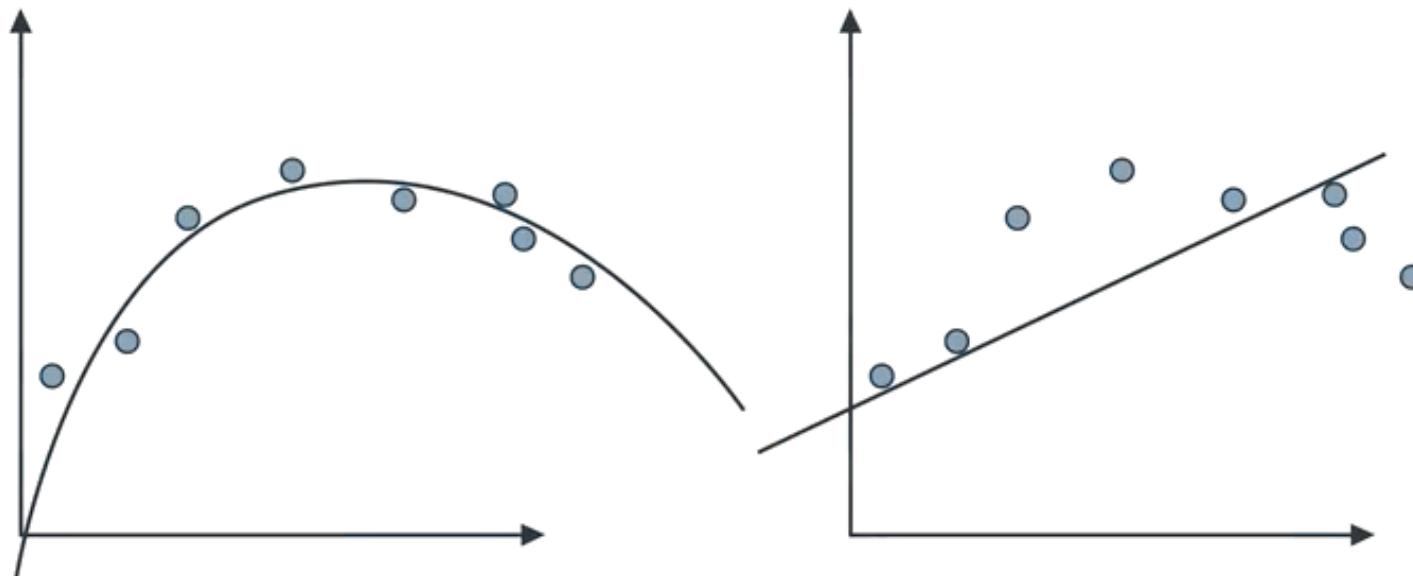


Error Detection – Underfitting

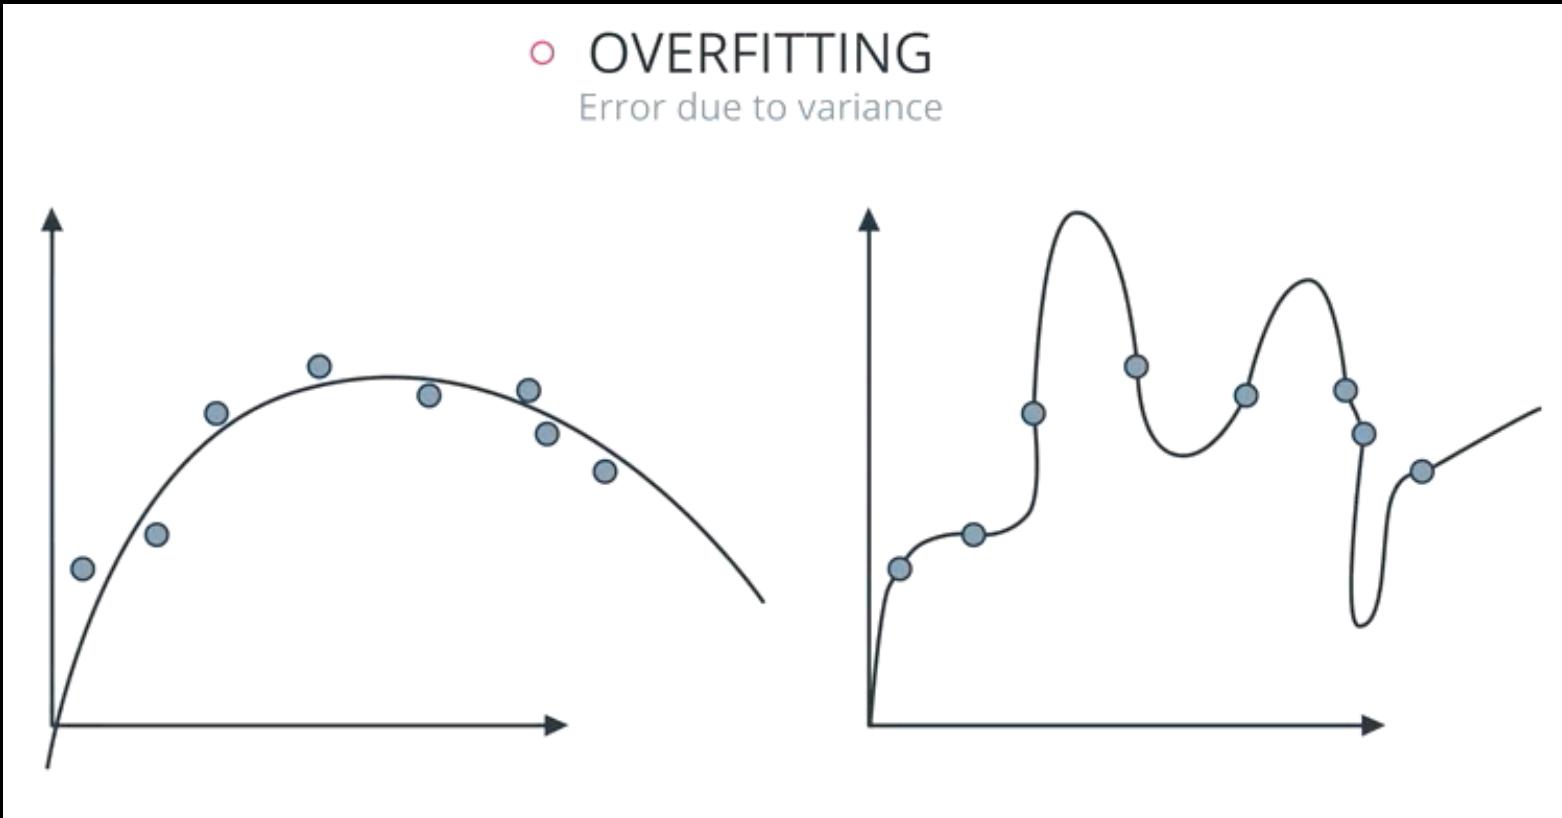
○ UNDERFITTING

Error due to bias

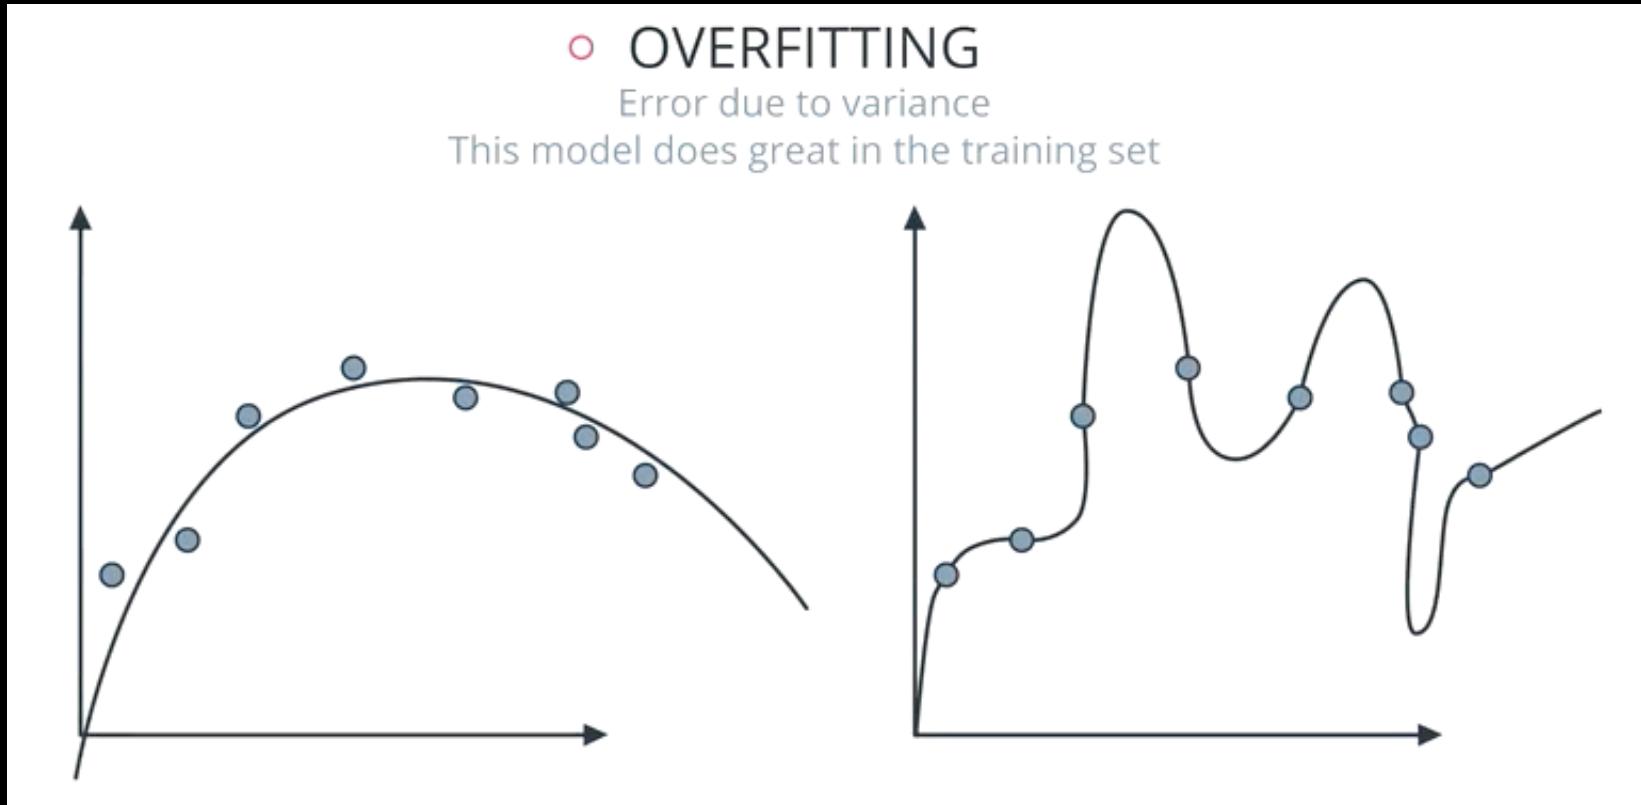
This model will not do well in the training set



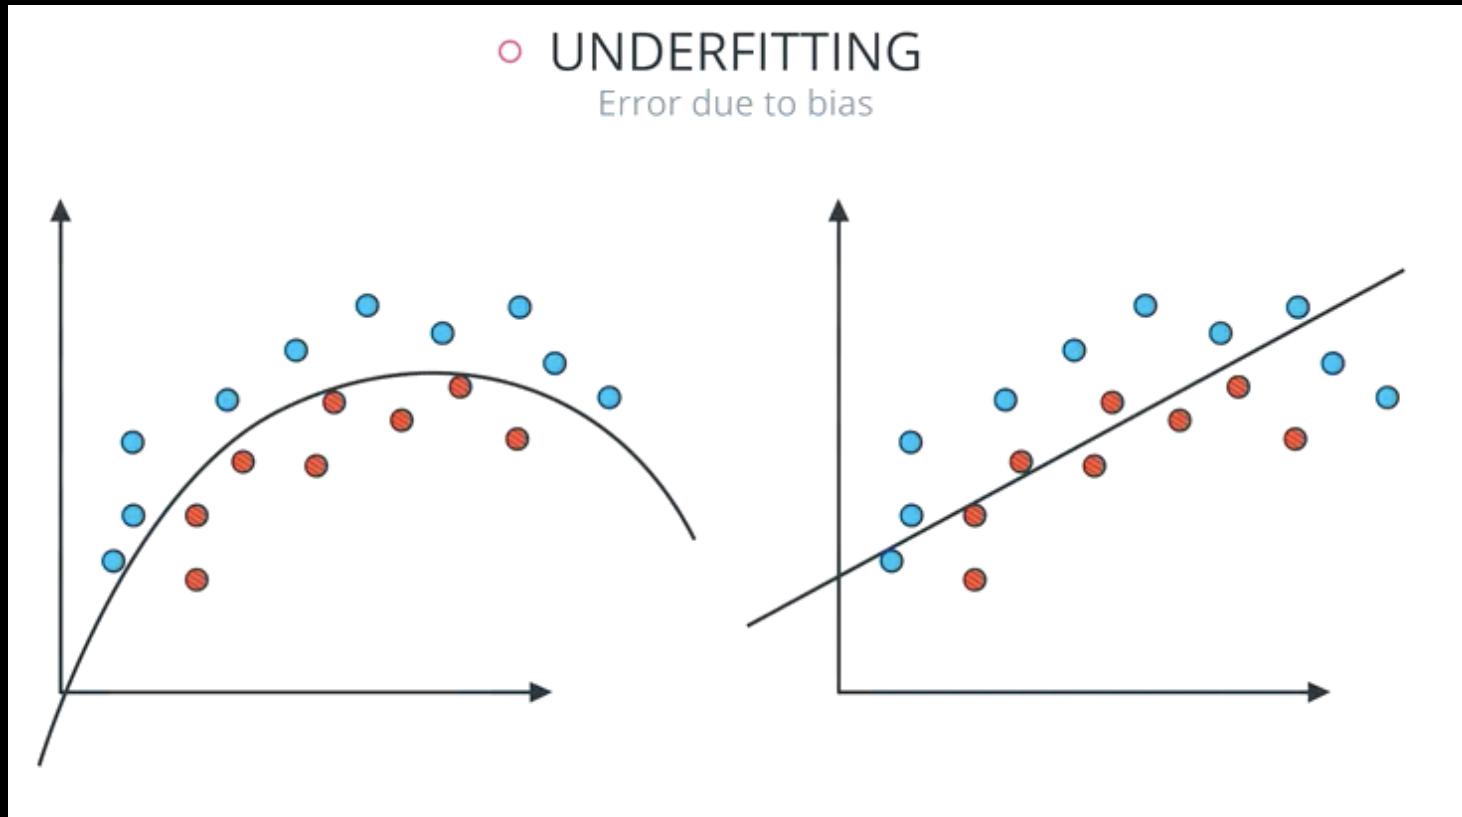
Error Detection – Overfitting



Error Detection – Overfitting



Error Detection – Underfitting



Error Detection – Overfitting



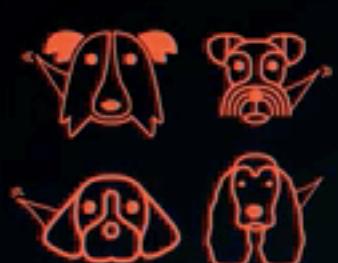
Error Detection – Summary

High bias (underfitting)

Not animals



Animals



Good Model

No dogs



Dogs



High variance (overfitting)

No dogs who wag their tails



Dogs who wag their tails



Oversimplify the problem

Bad on training set

Bad on testing set

Good model

Good on training set

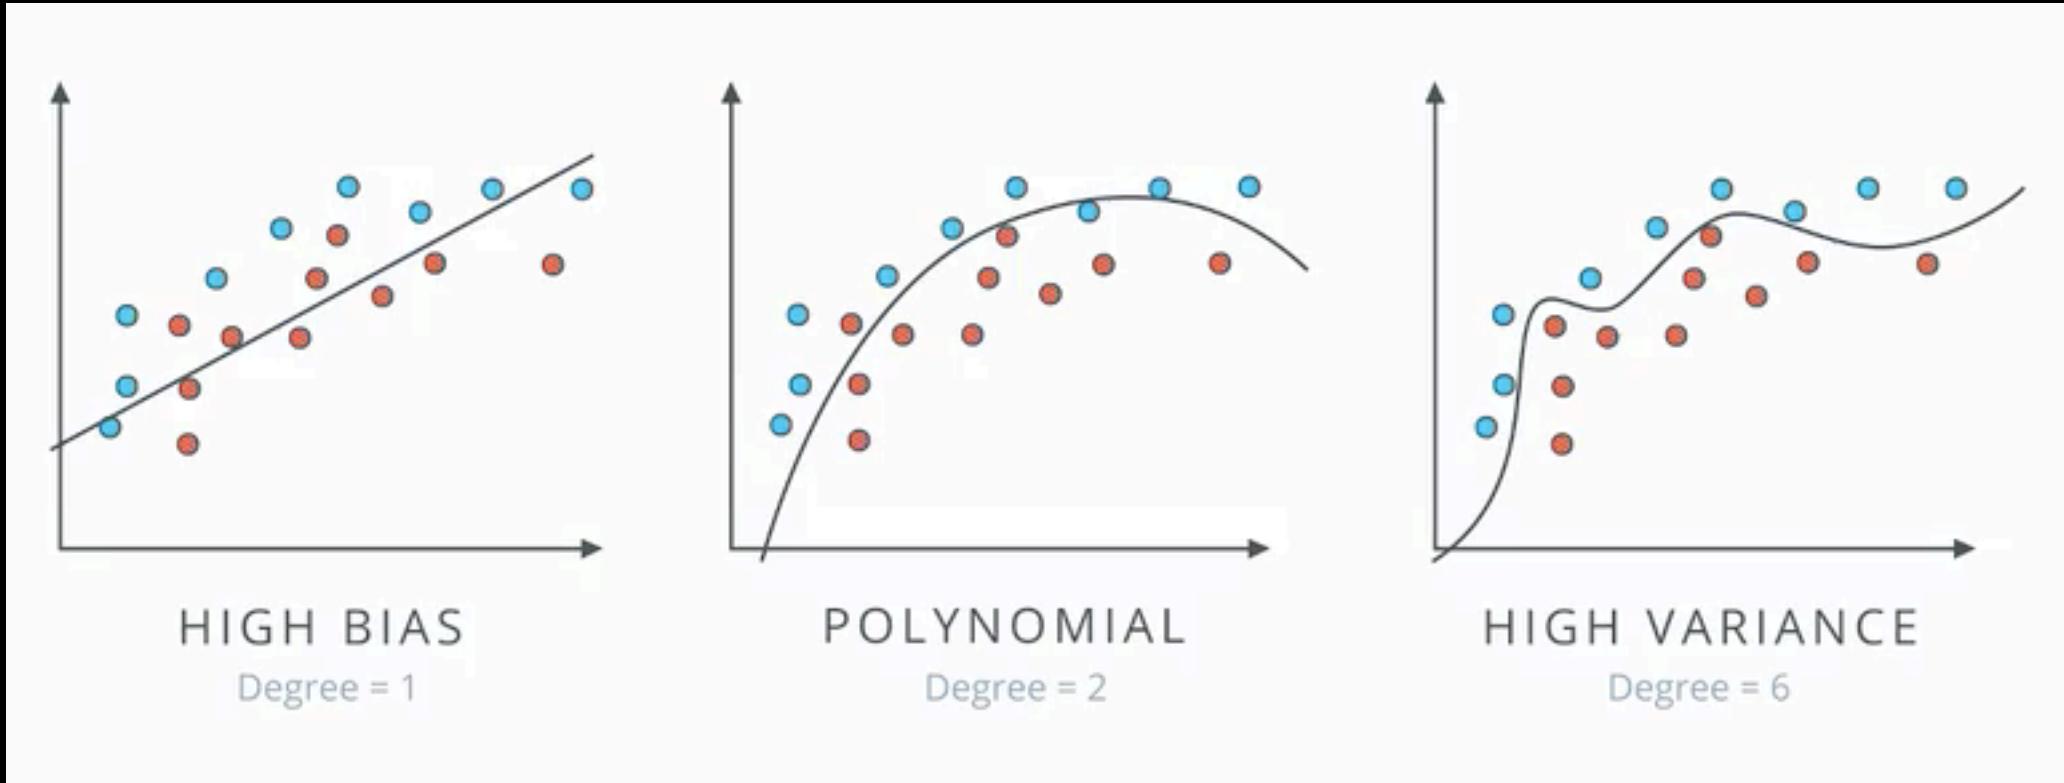
Good on testing set

Overcomplicate the problem

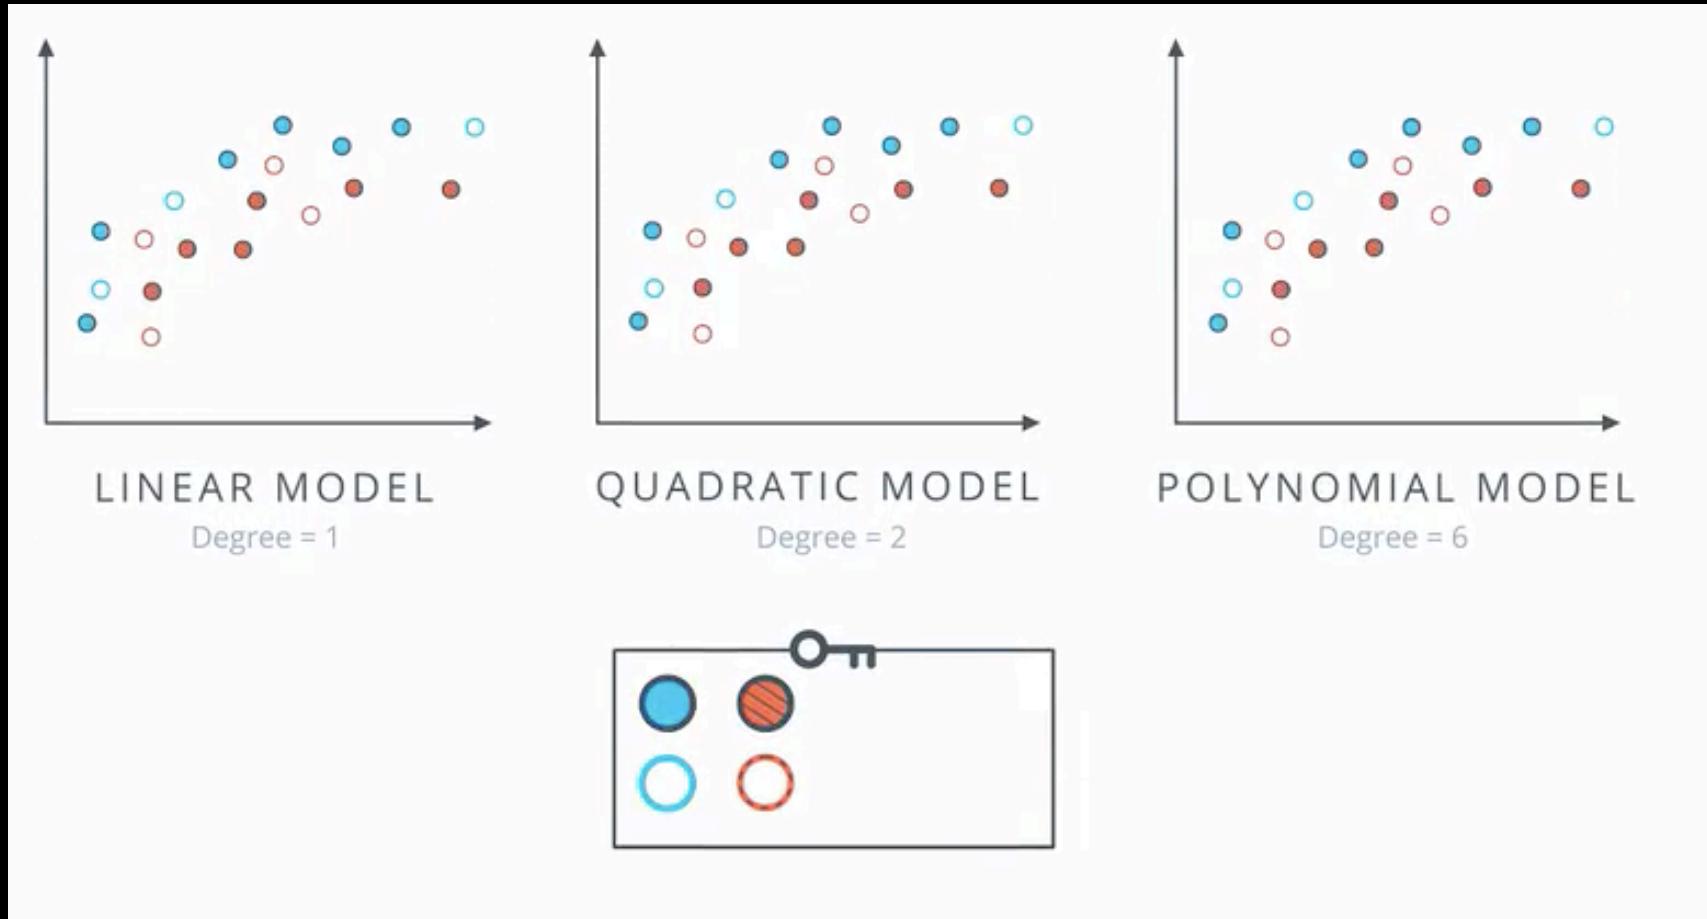
Great on training set

Bad on testing set

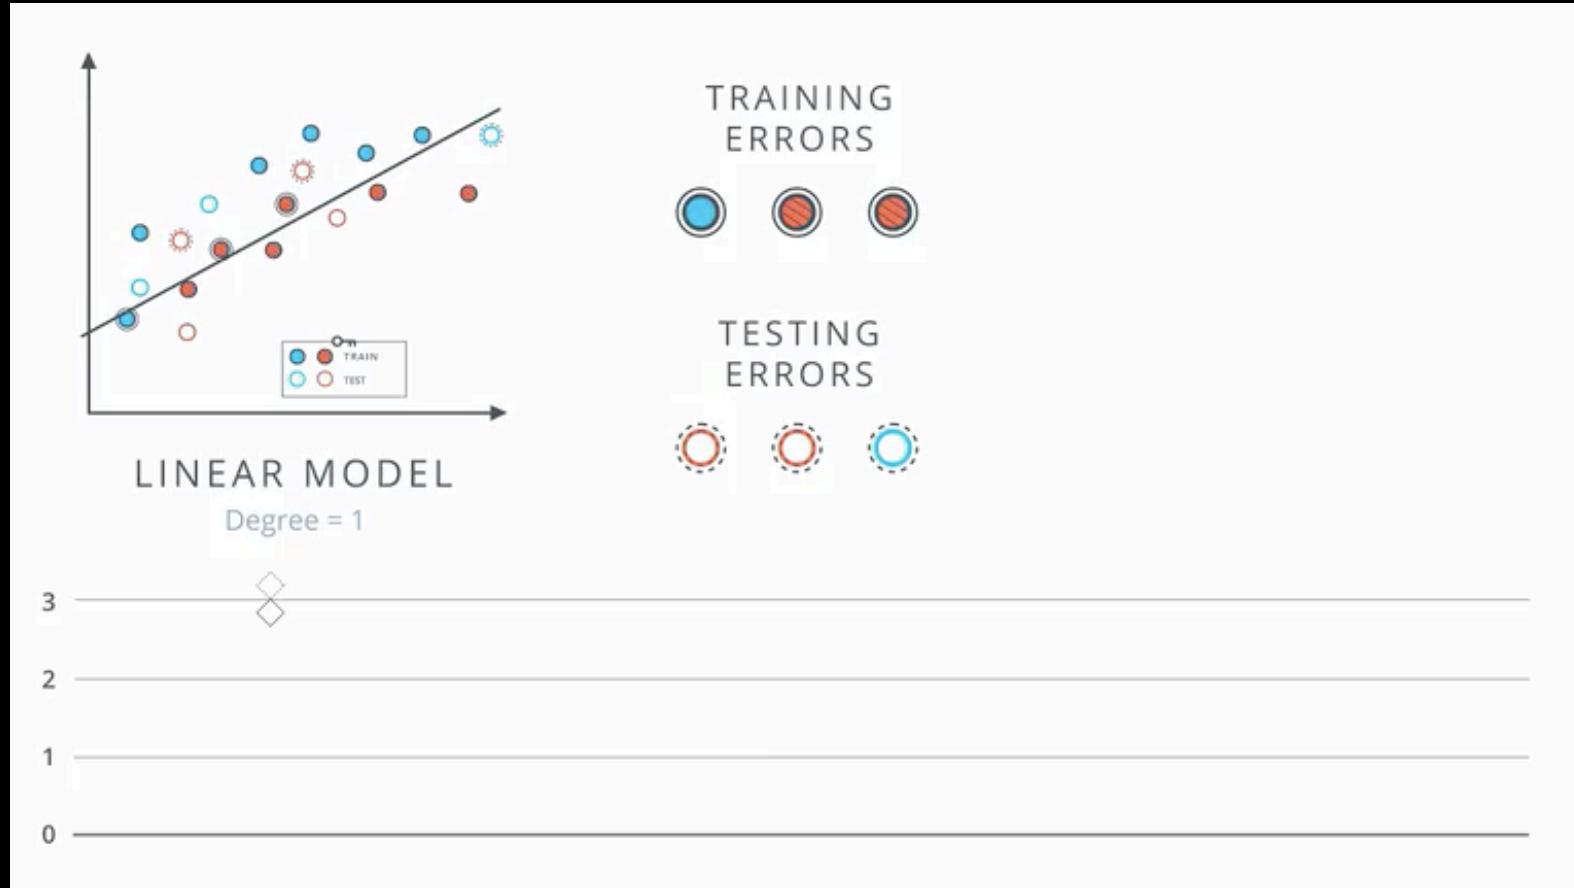
Model Complexity Graph



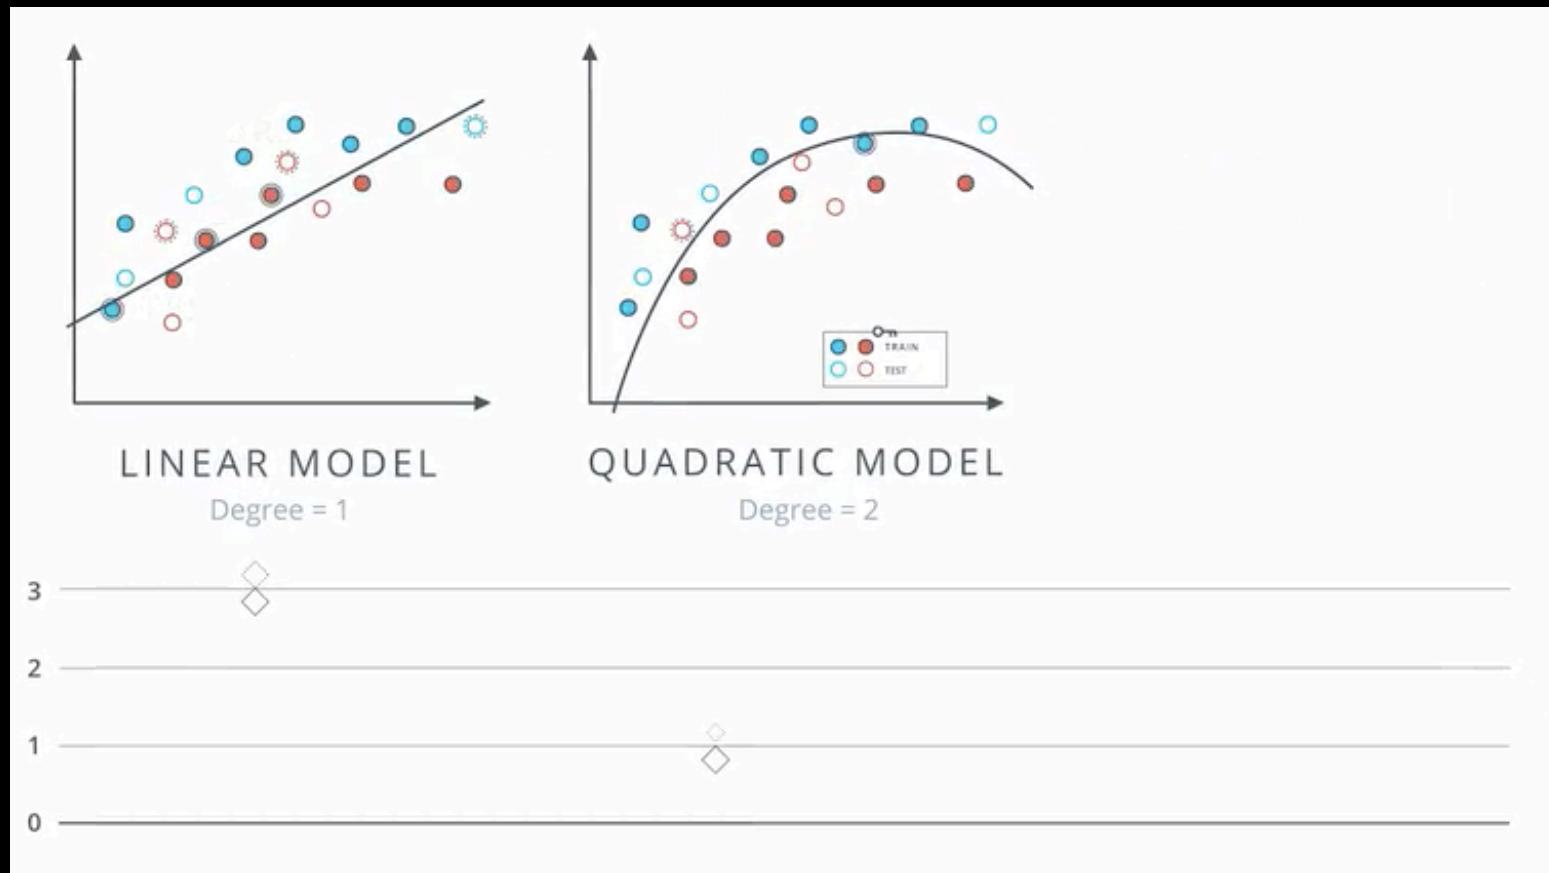
Model Complexity Graph



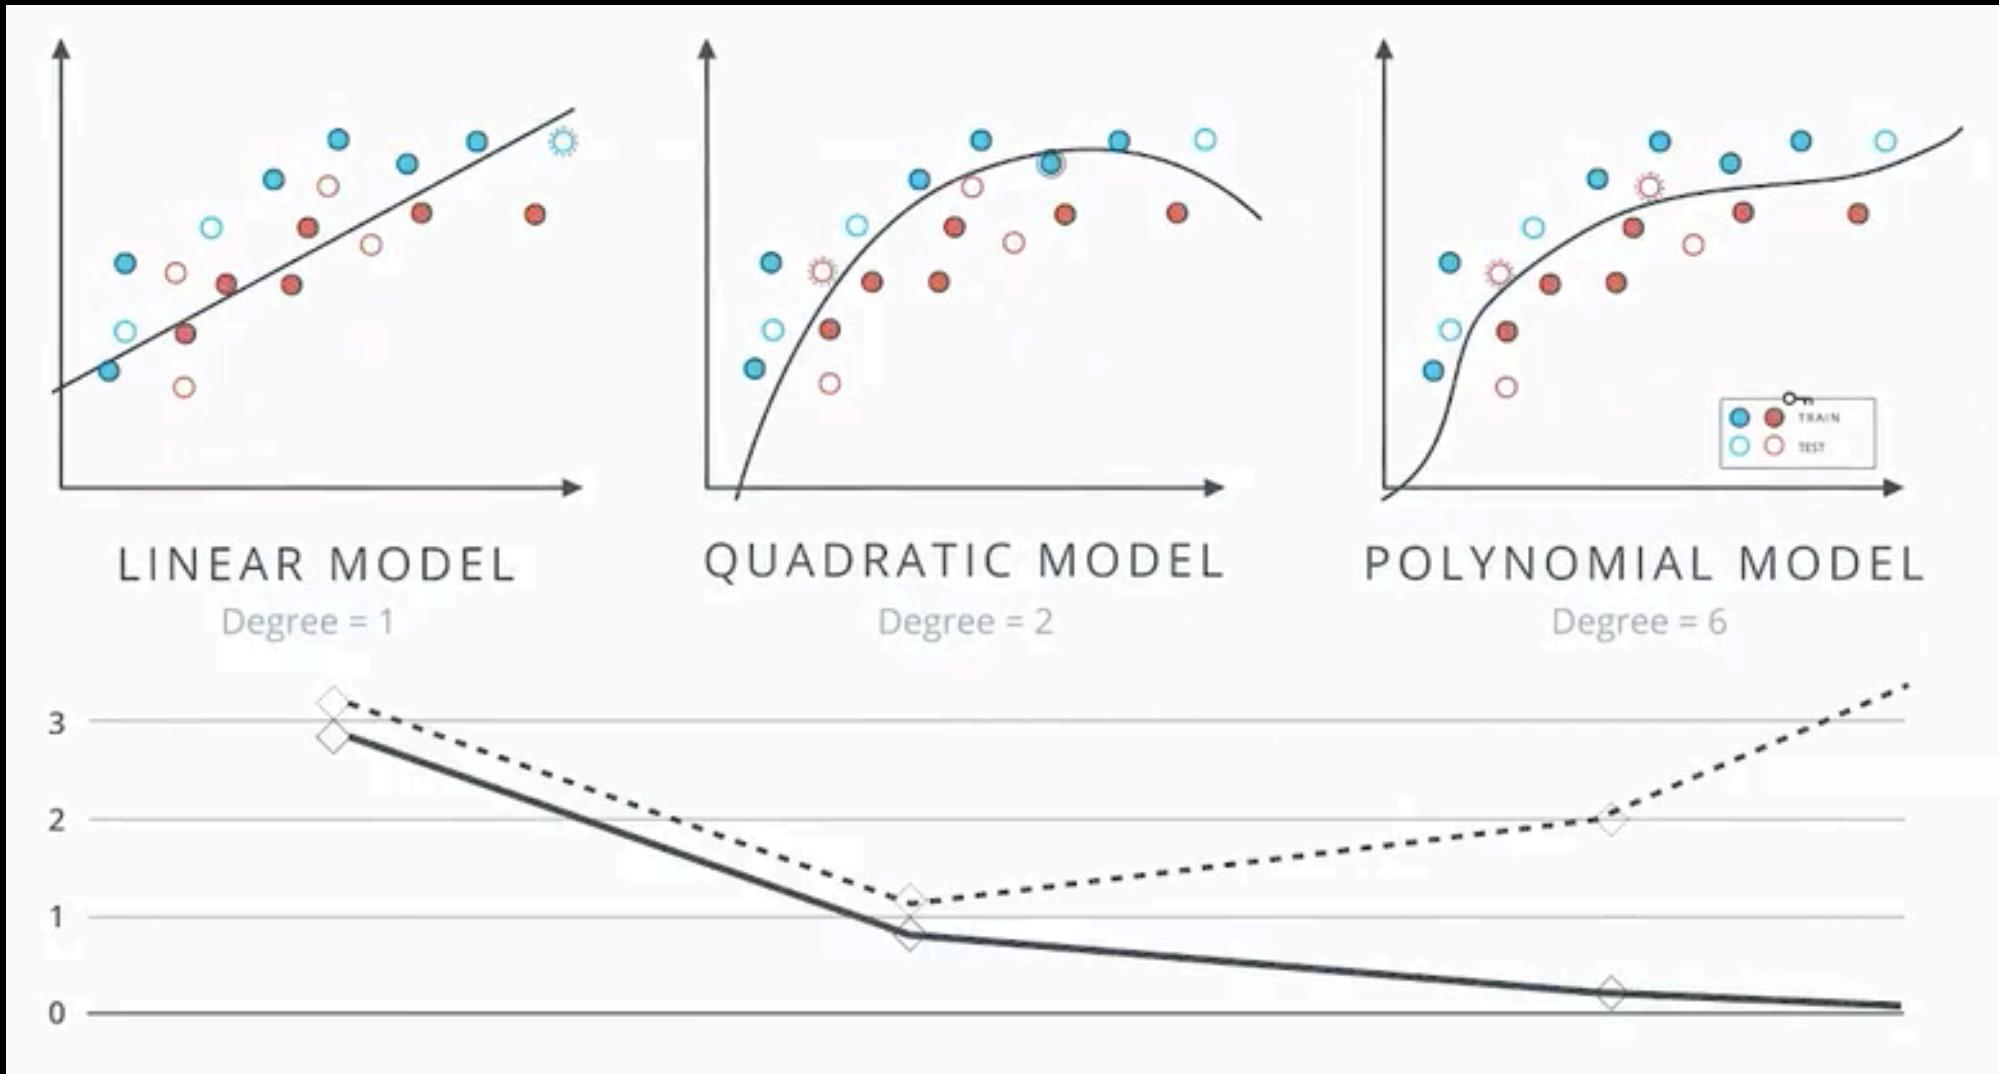
Model Complexity Graph



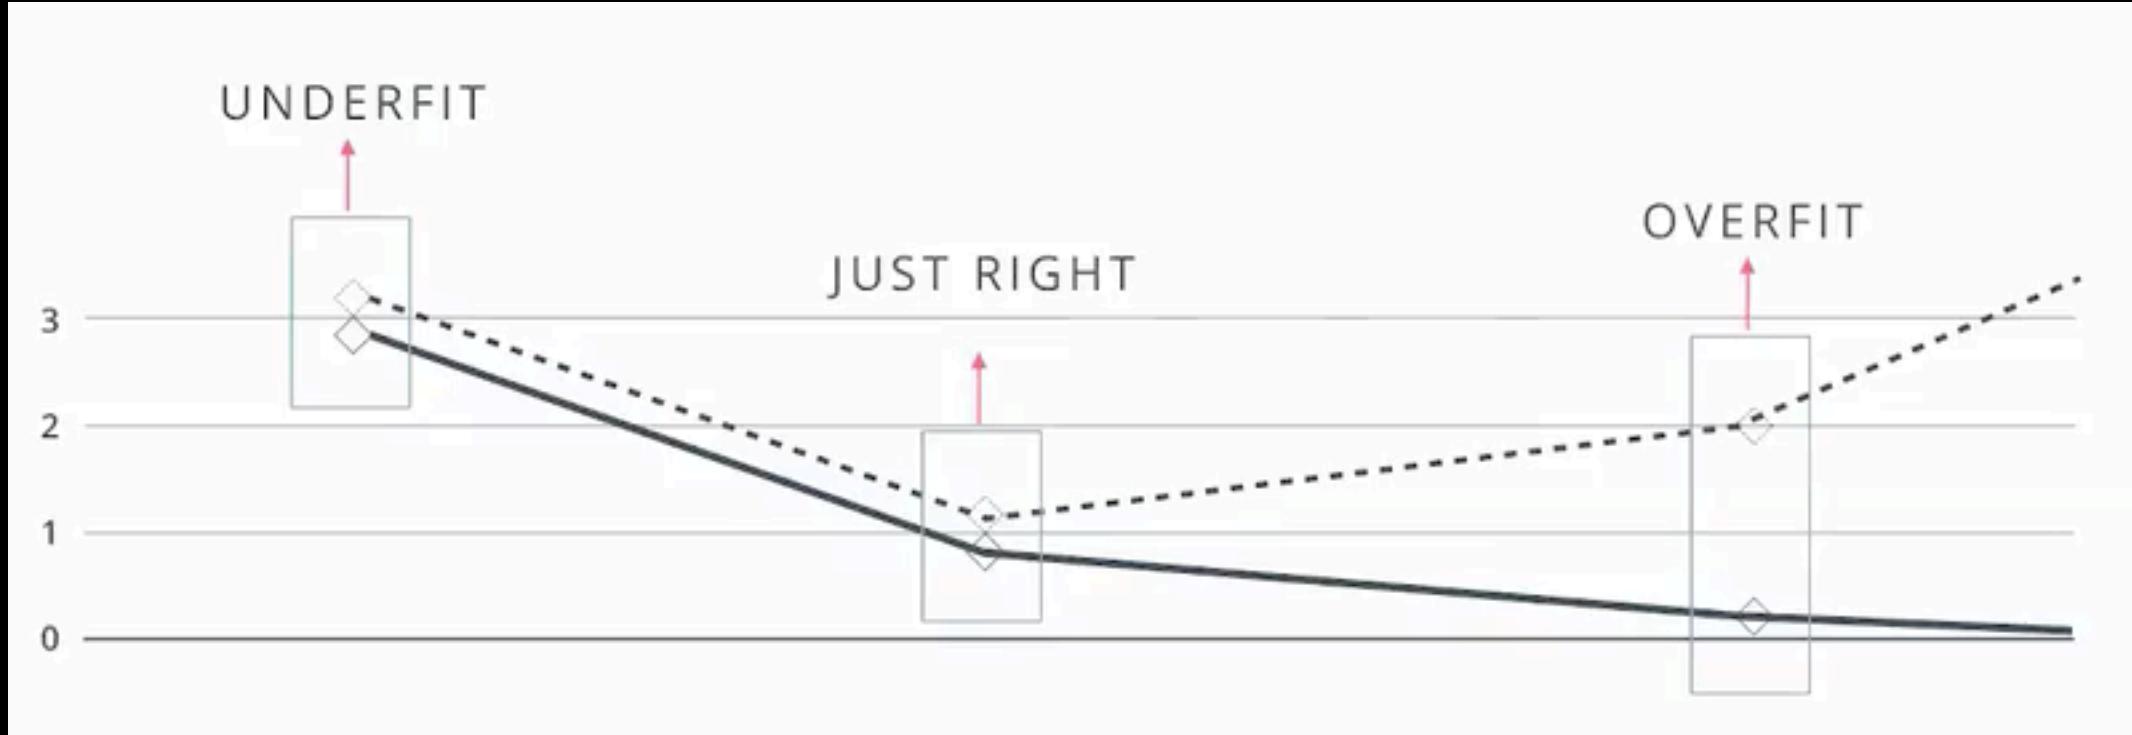
Model Complexity Graph



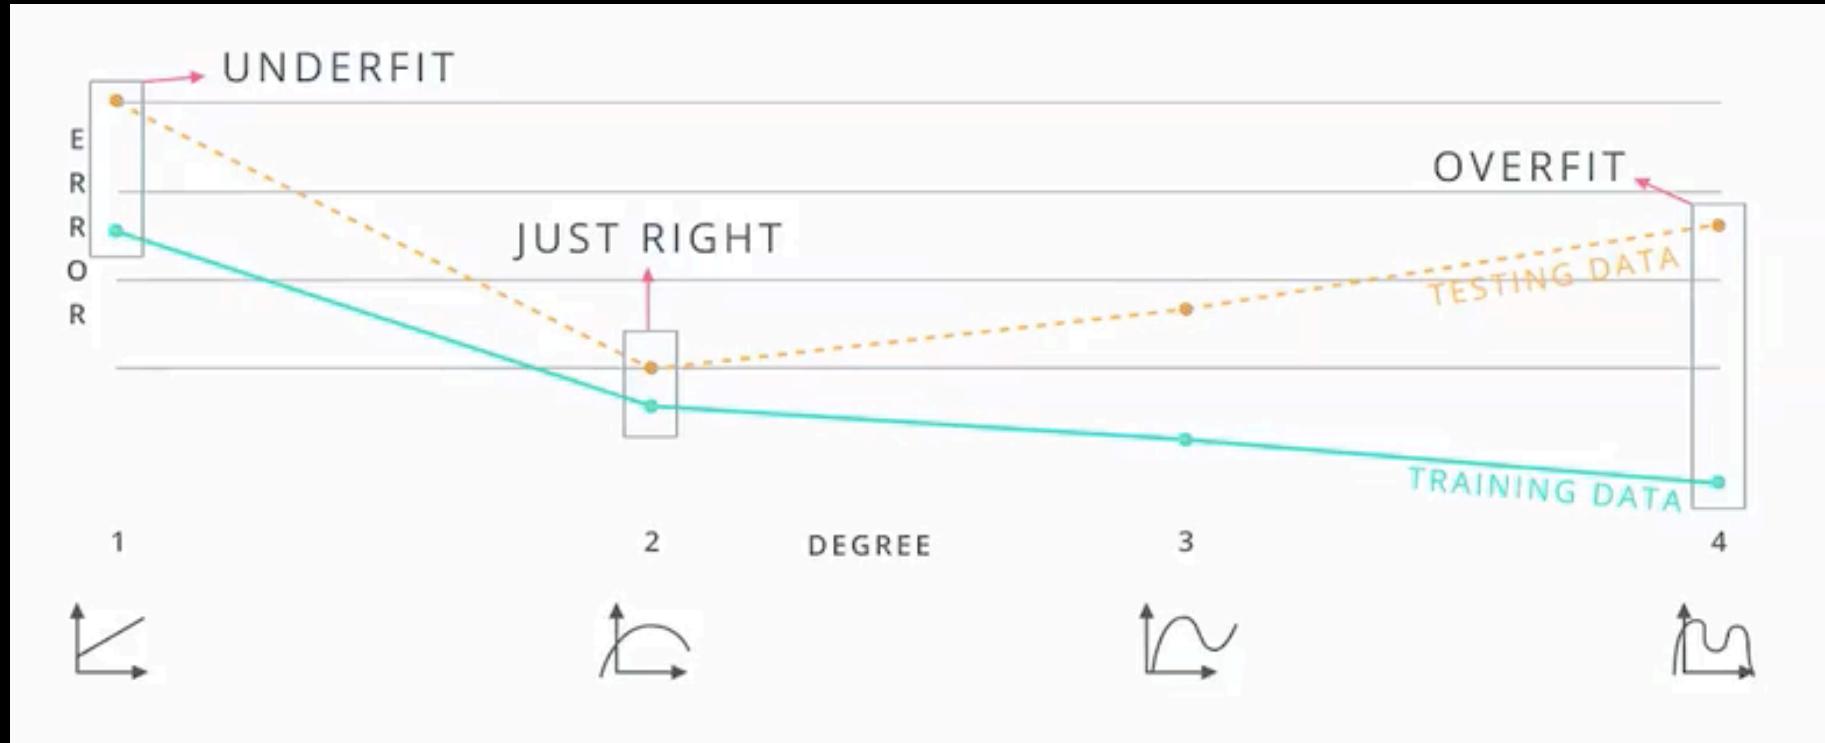
Model Complexity Graph



Model Complexity Graph



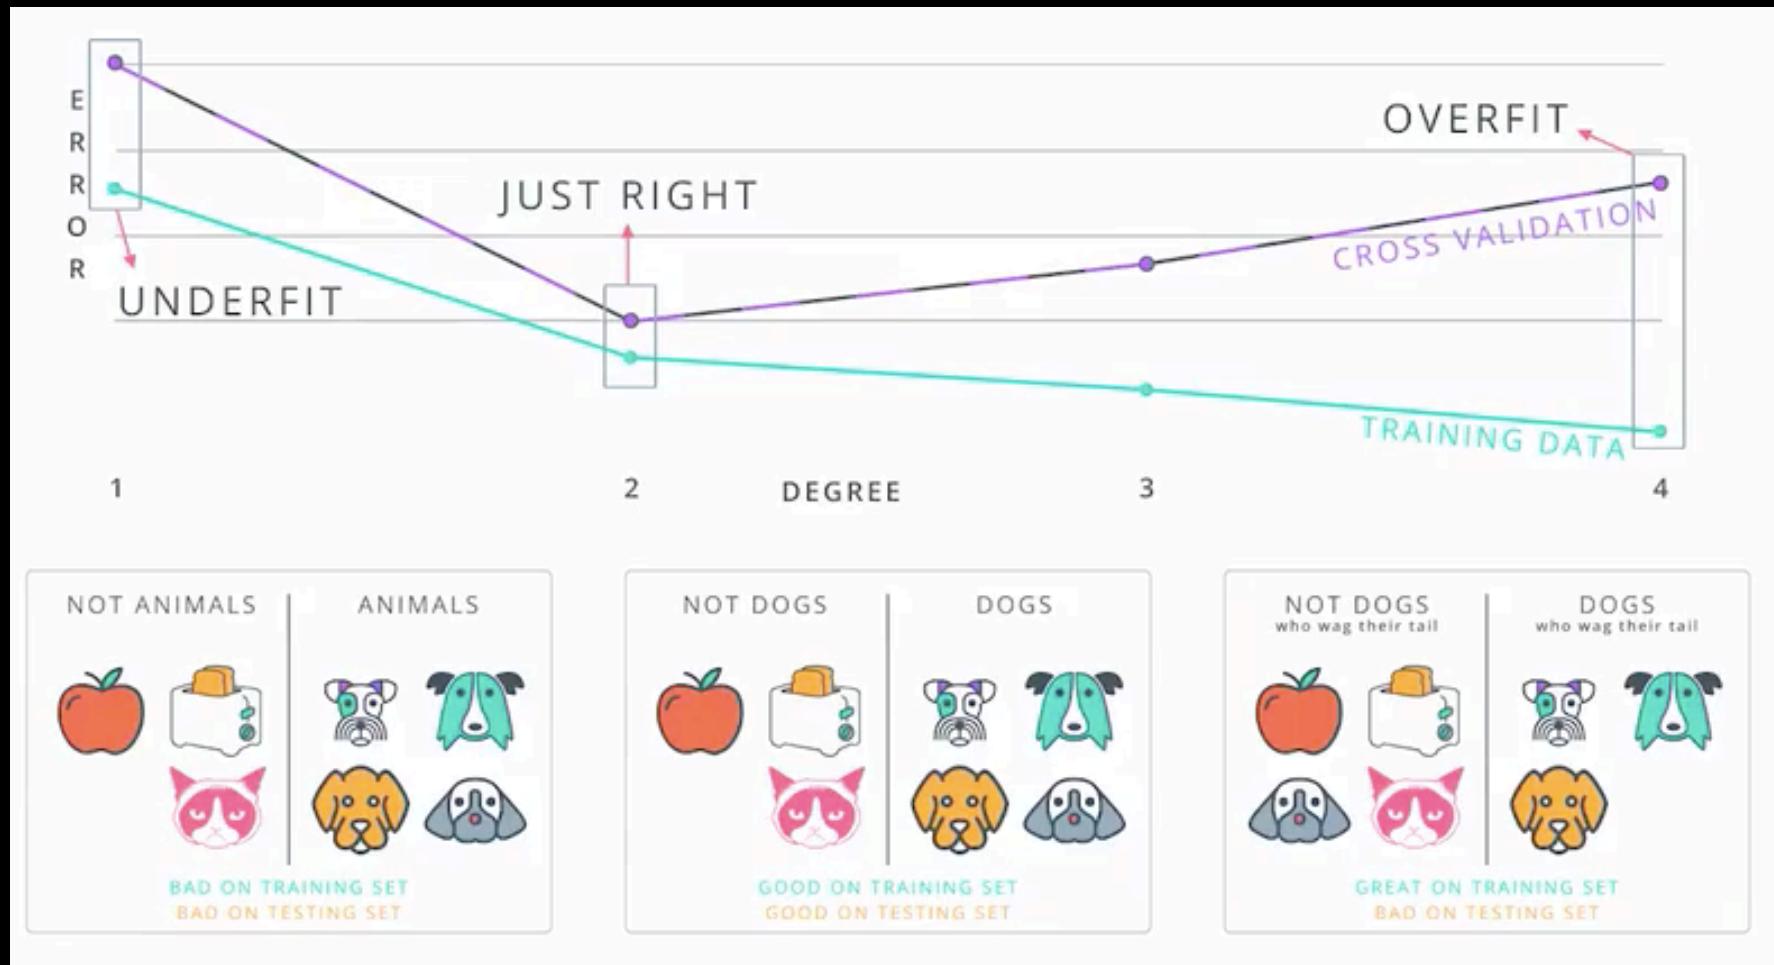
Model Complexity Graph



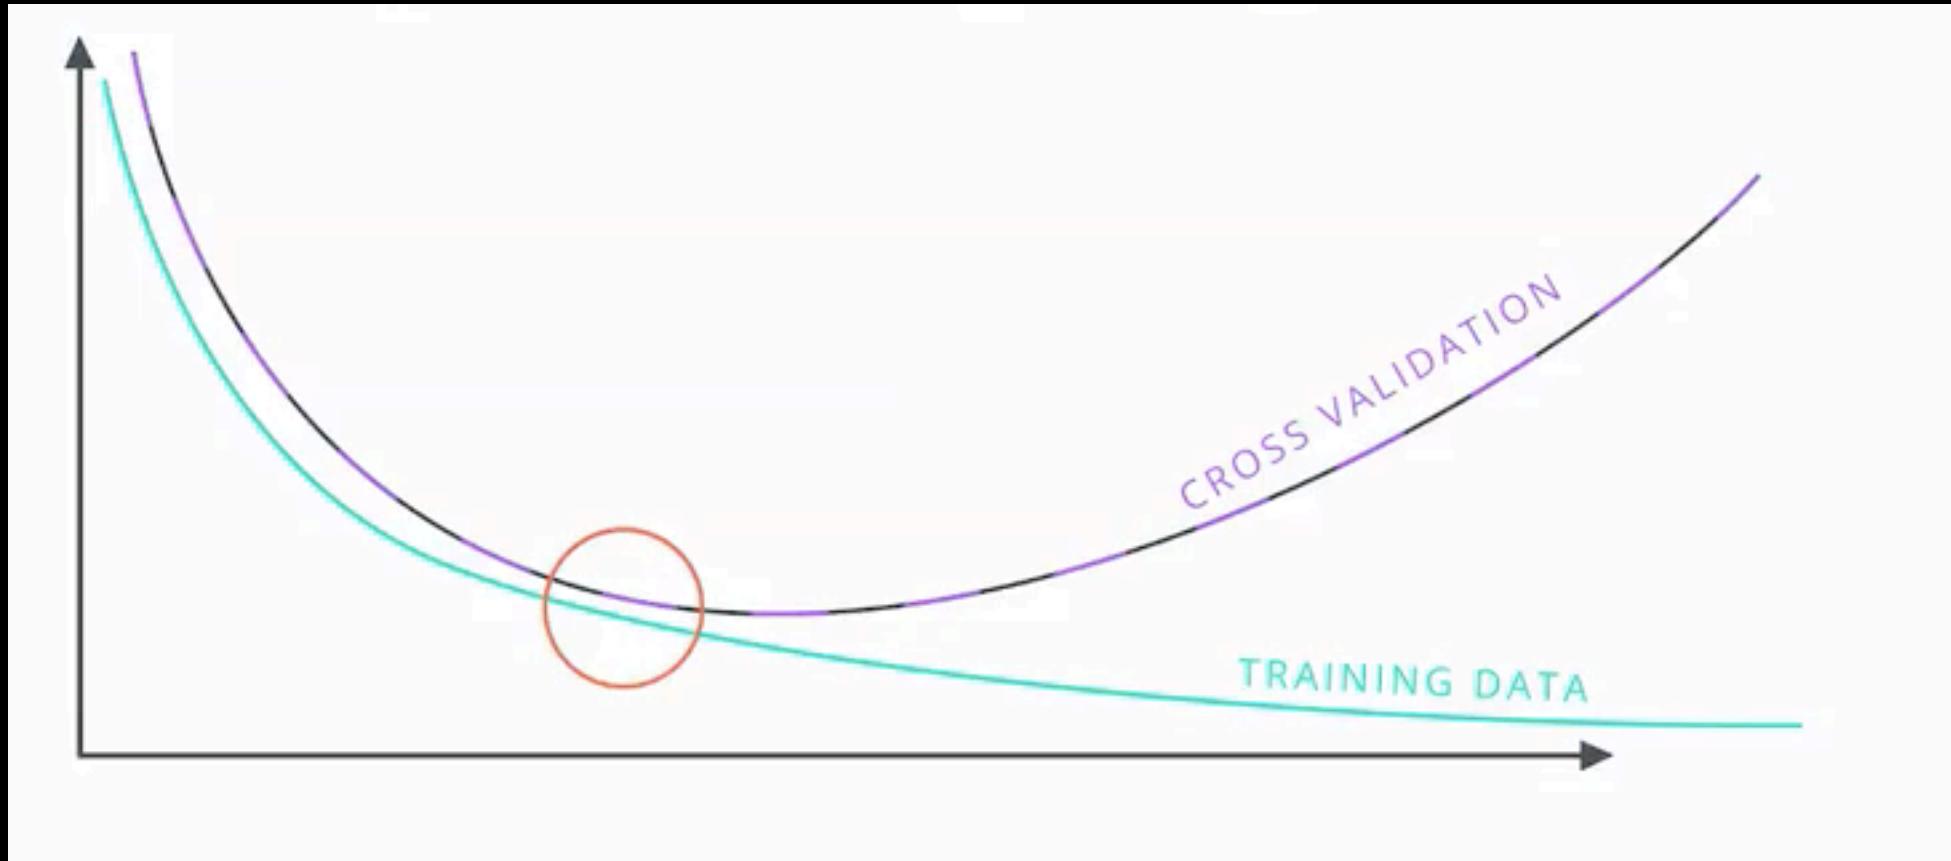
Cross Validation



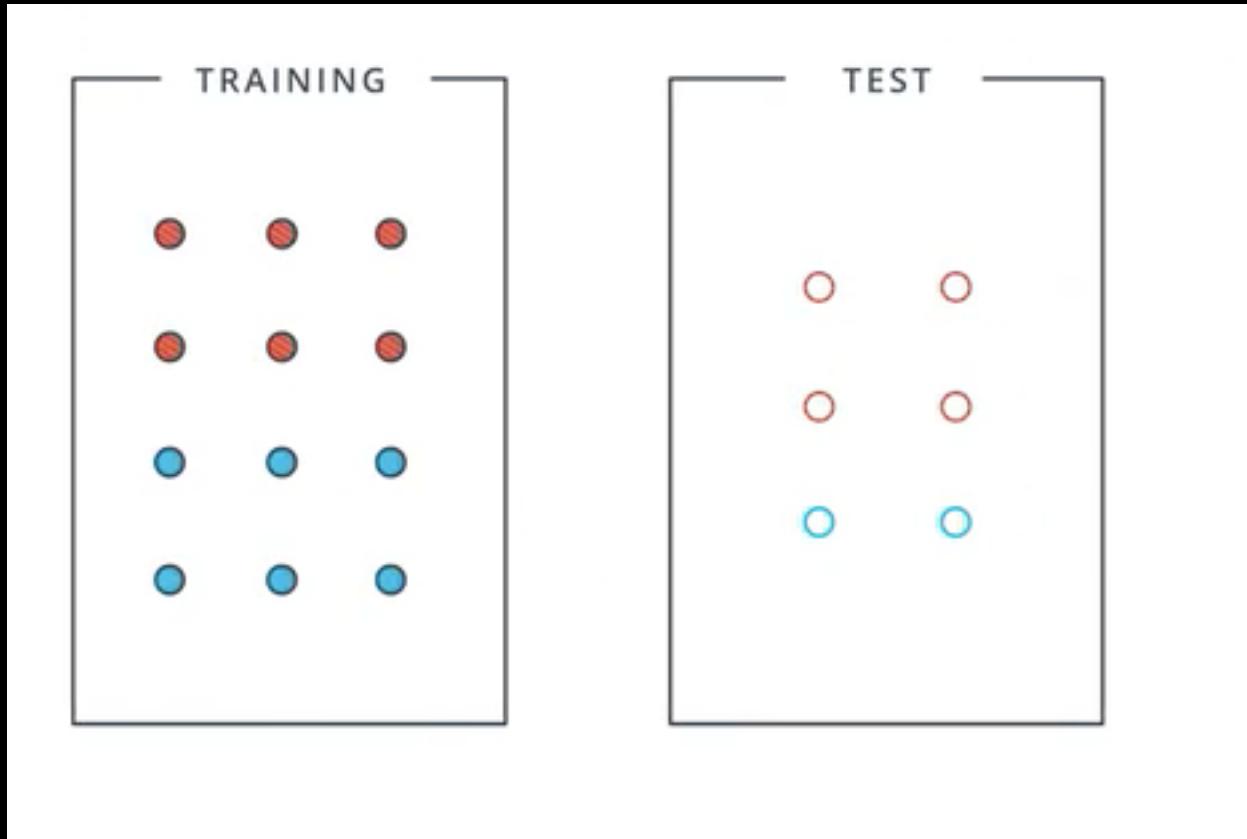
Cross Validation



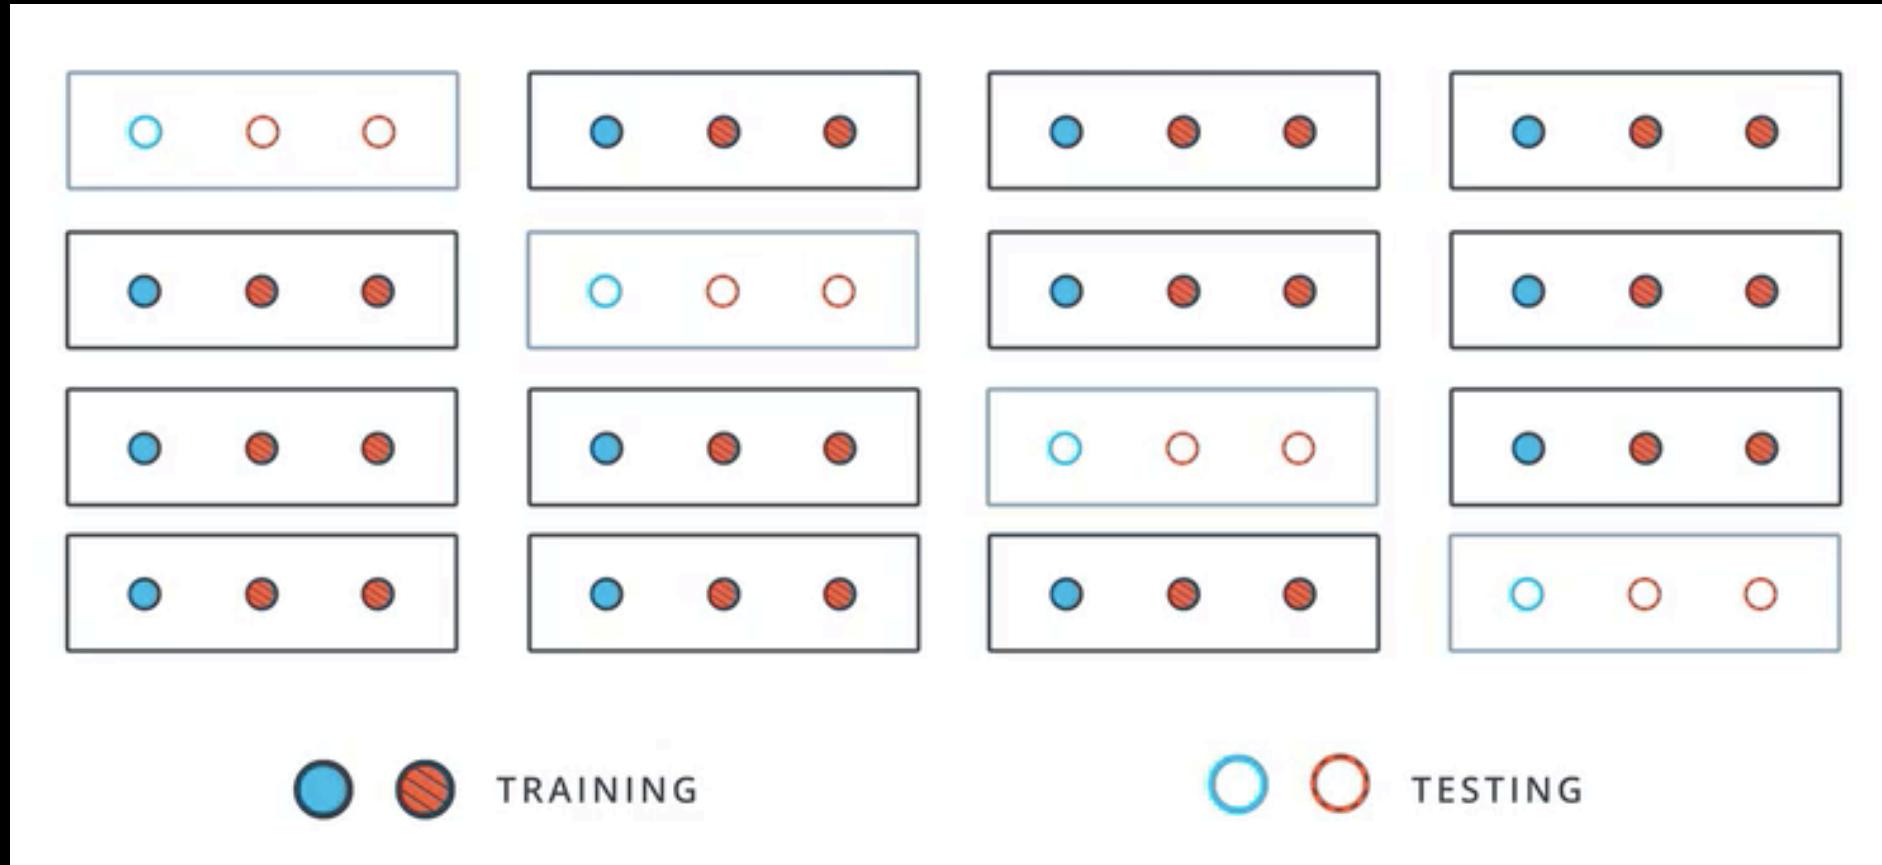
Cross Validation



K-Fold Cross Validation



K-Fold Cross Validation

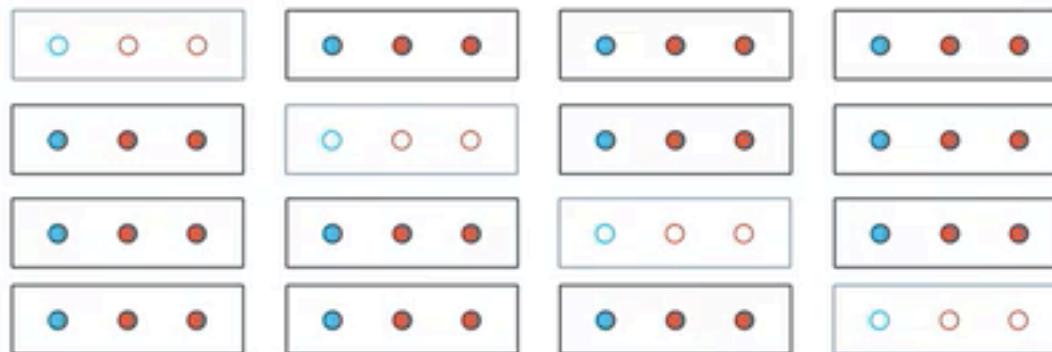


K-Fold Cross Validation

```
from sklearn.model_selection import KFold  
kf = KFold(12, 3)
```

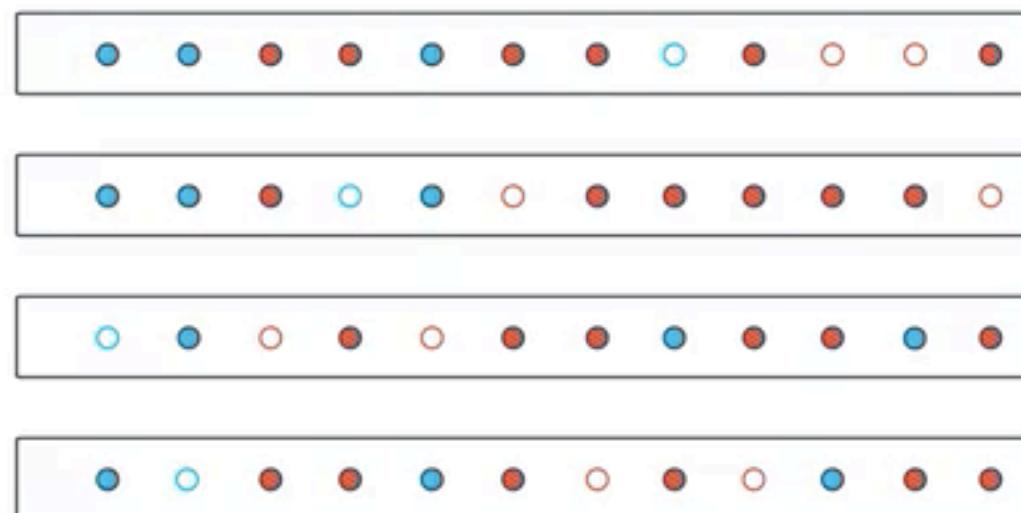
```
for train_indices, test_indices in kf:  
    print train_indices, test_indices
```

```
[3 4 4 6 7 8 9 10 11] [0 1 2]  
[0 1 2 6 7 8 9 10 11] [3 4 5]  
[0 1 2 3 4 5 9 10 11] [6 7 8]  
[0 1 2 3 4 5 6 4 8 ] [9 10 11]
```

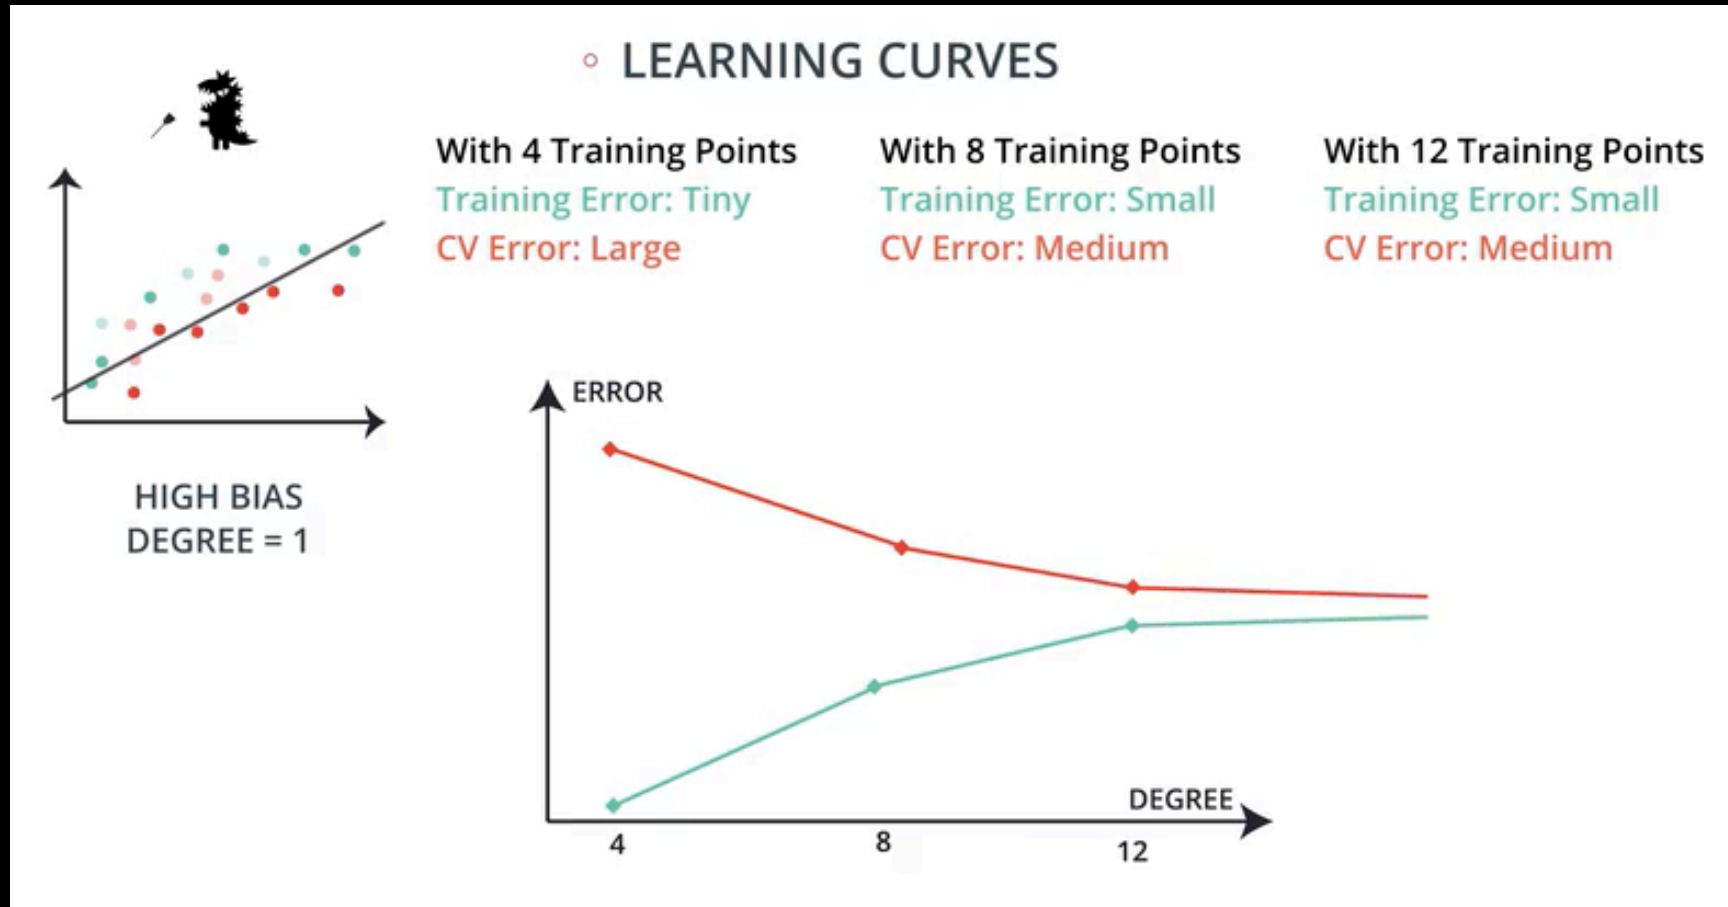


K-Fold Cross Validation

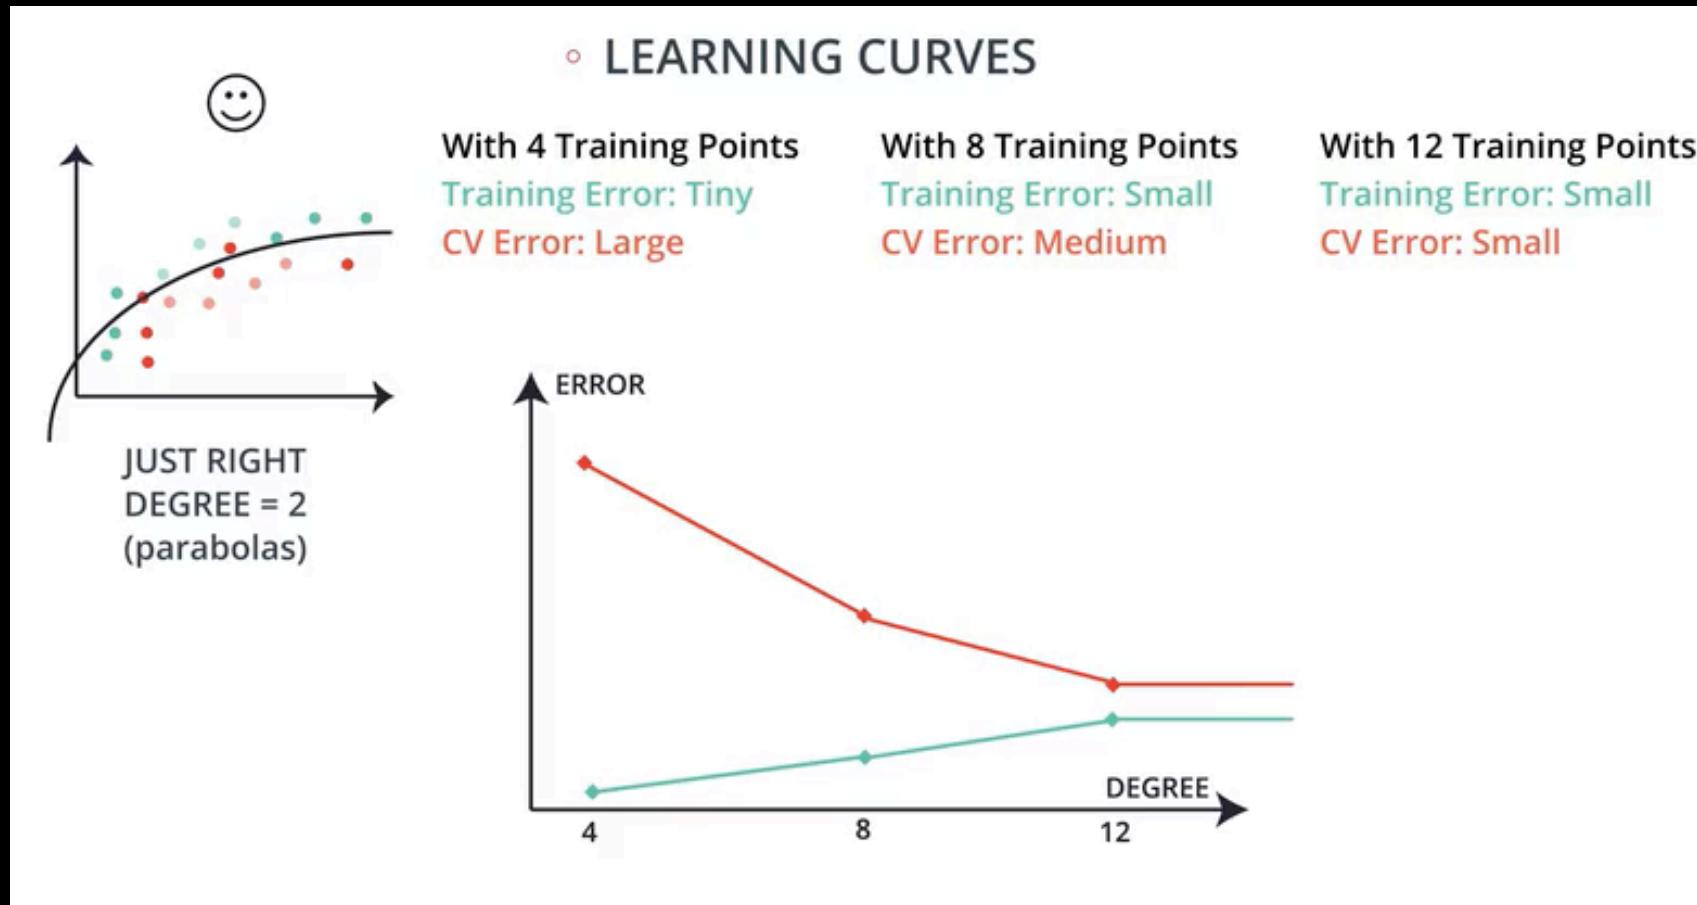
```
from sklearn.model_selection import KFold  
kf = KFold(12, 3, shuffle = True)  
  
for train_indices, test_indices in kf:  
    print train_indices, test_indices  
  
[0 1 2 3 4 5 6 8 11] [7 9 10]  
[0 1 2 4 6 7 8 9 10] [3 5 11]  
[1 3 5 6 7 8 9 10 11] [0 2 4]  
[0 2 3 4 5 7 9 10 11] [1 6 8]
```



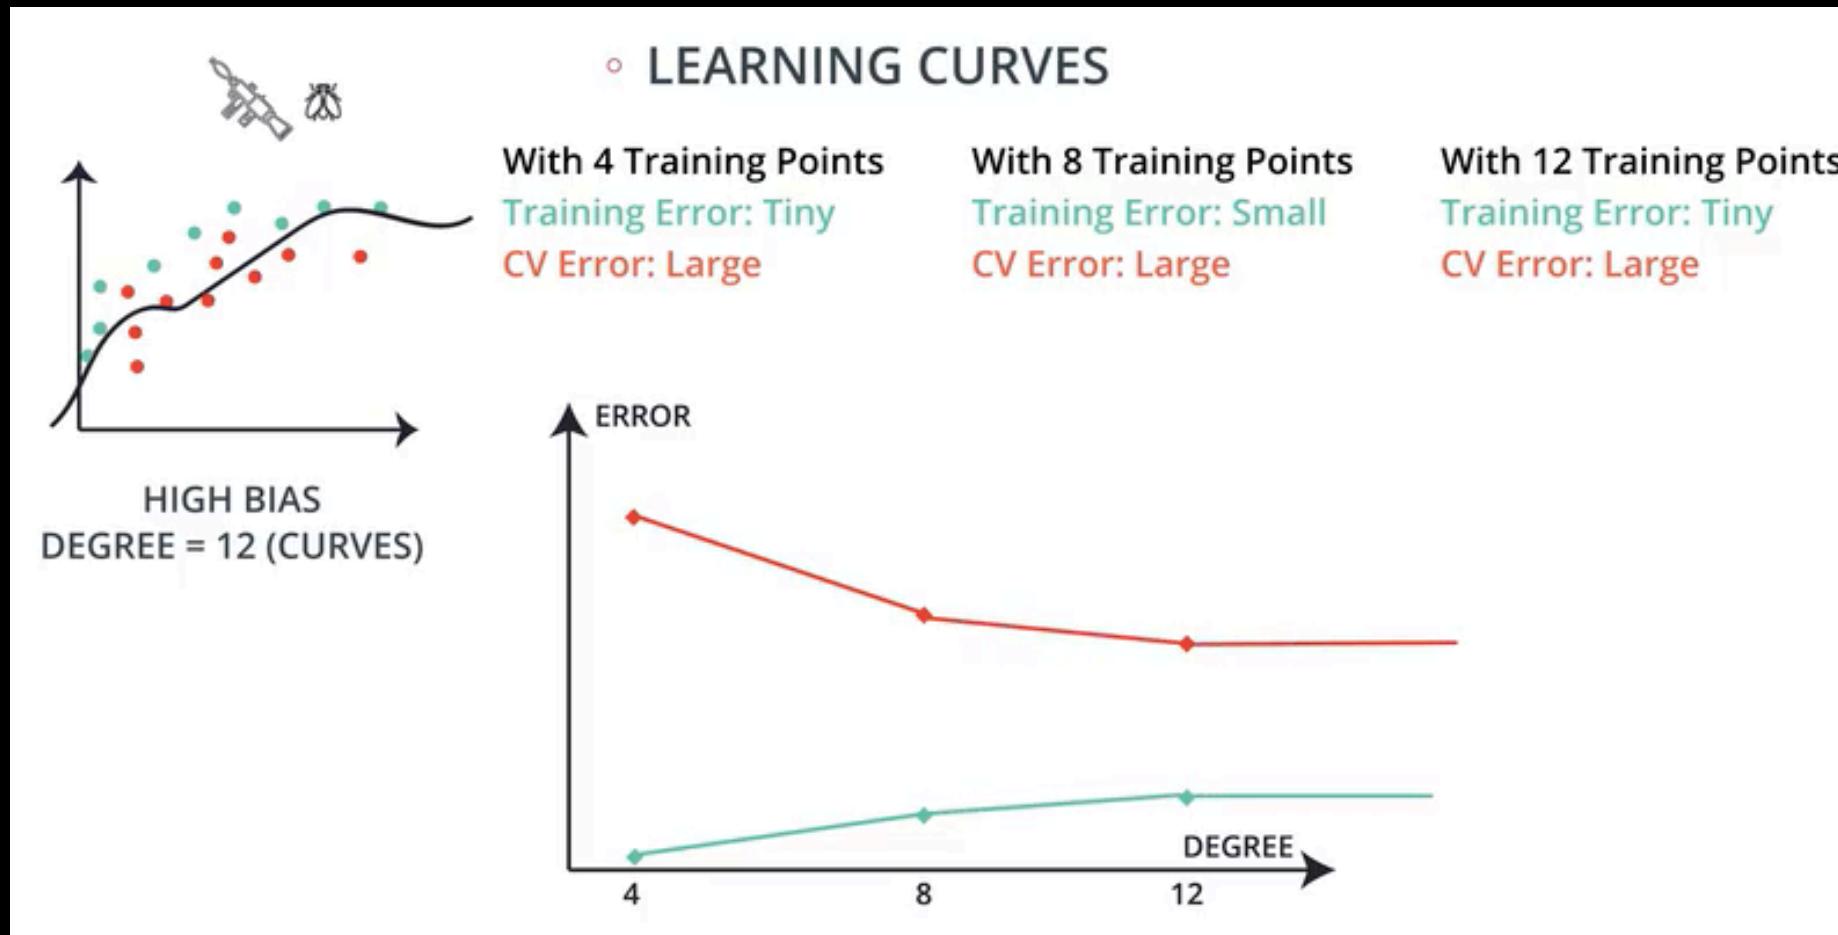
Learning Curves Summary



Learning Curves Summary



Learning Curves Summary



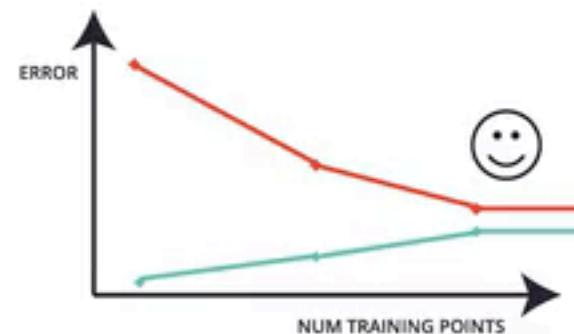
Learning Curves Summary

◦ LEARNING CURVES

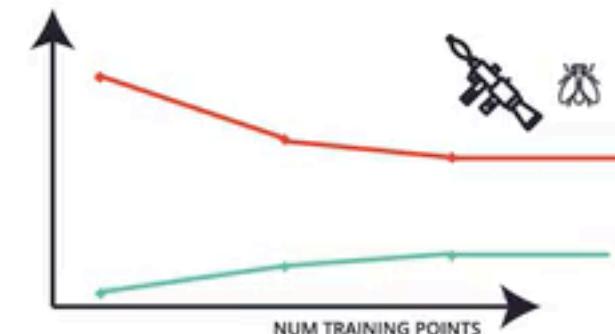
TRAINING ERROR
CV ERROR



HIGH BIAS



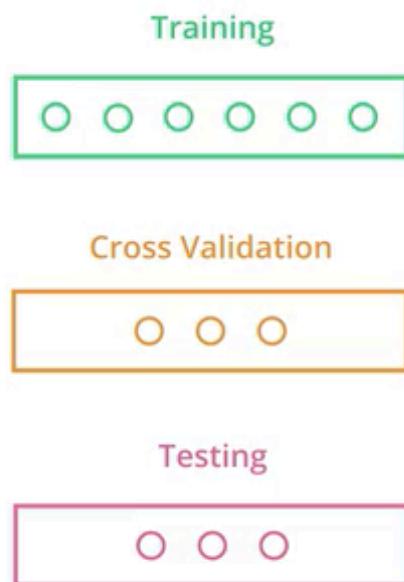
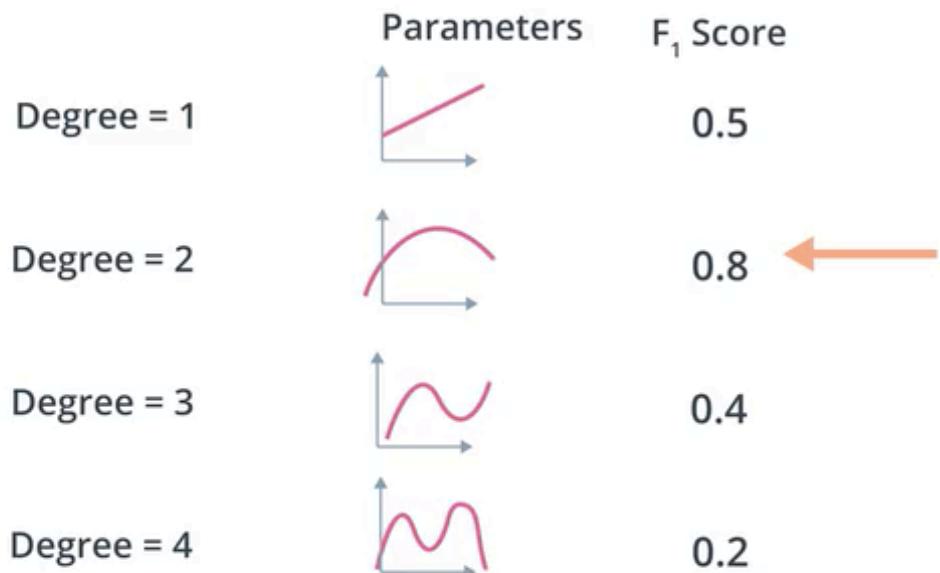
JUST RIGHT



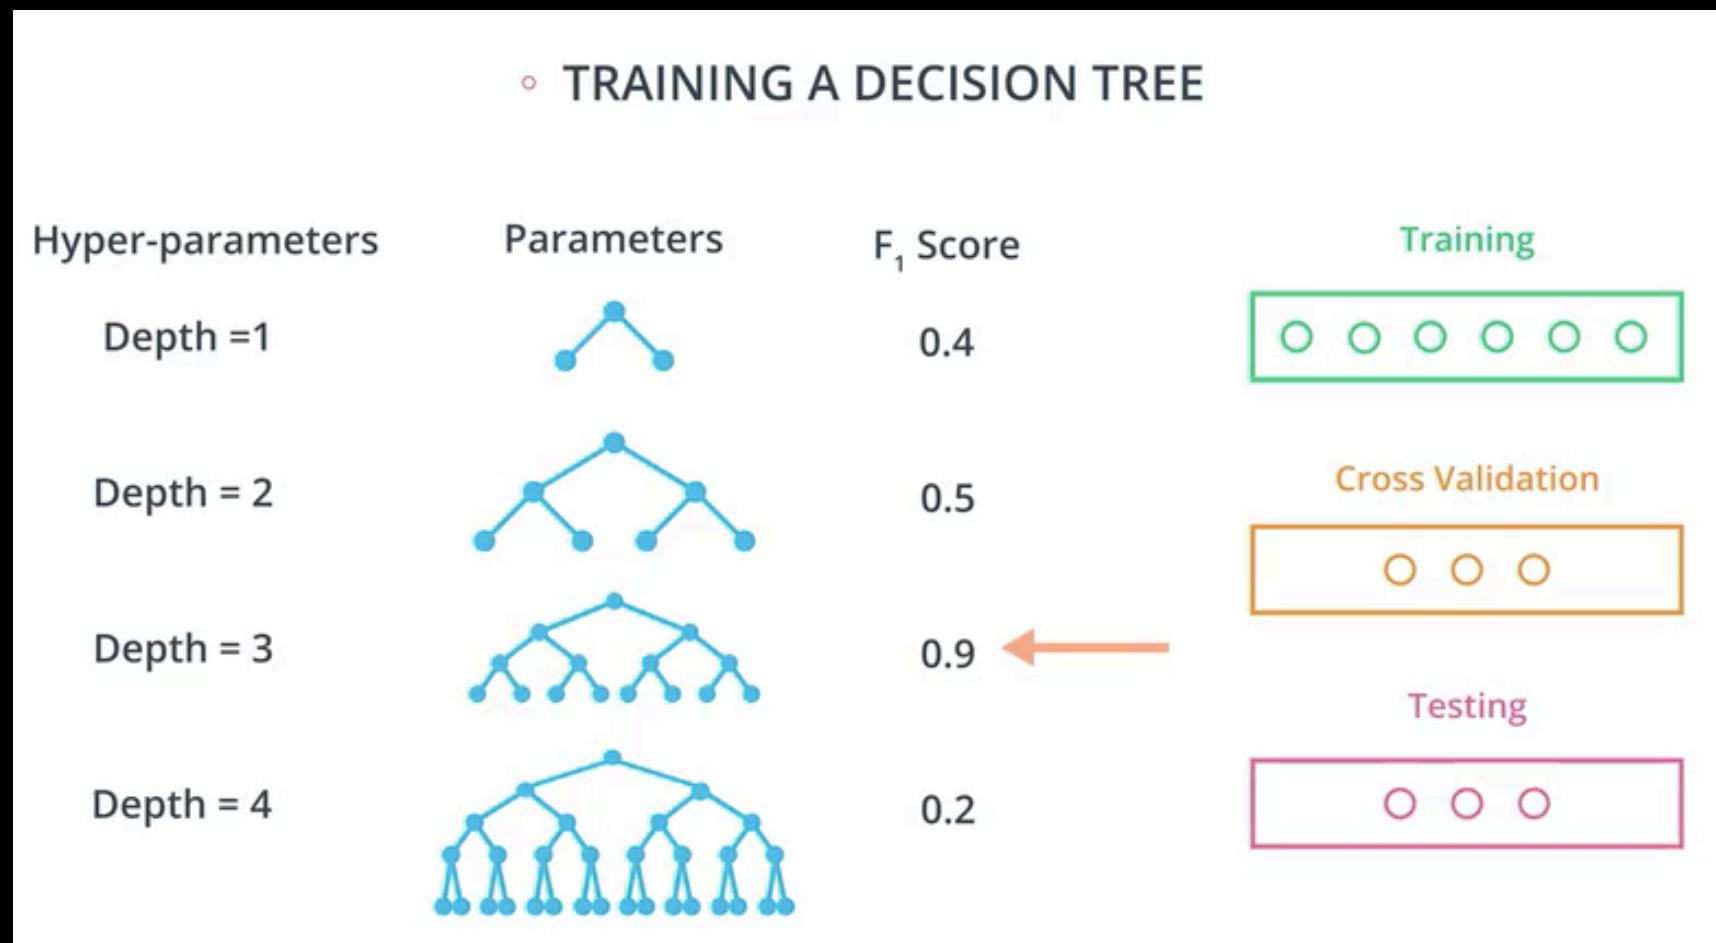
HIGH VARIANCE

Grid Search

- TRAINING A LOGISTIC REGRESSION MODEL



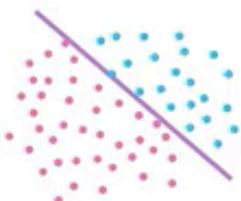
Grid Search



Grid Search

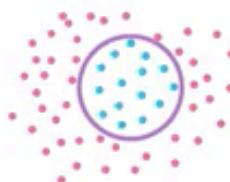
- TRAINING A SUPPORT VECTOR MACHINE

Hyper-parameters



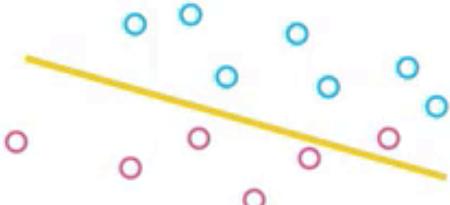
Kernel

Linear

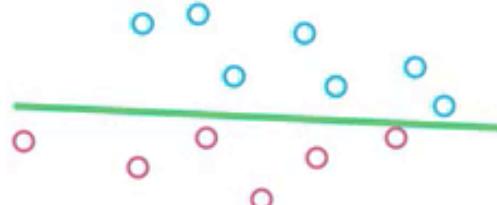


Polynomial

C



Small



Large

Grid Search

- GRID SEARCH CROSS VALIDATION

Kernel C	Linear	Polynomial
0.1	 F1 SCORE = 0.5	 F1 SCORE = 0.2
1	 F1 SCORE = 0.8	 F1 SCORE = 0.4
10	 F1 SCORE = 0.6	 F1 SCORE = 0.6

